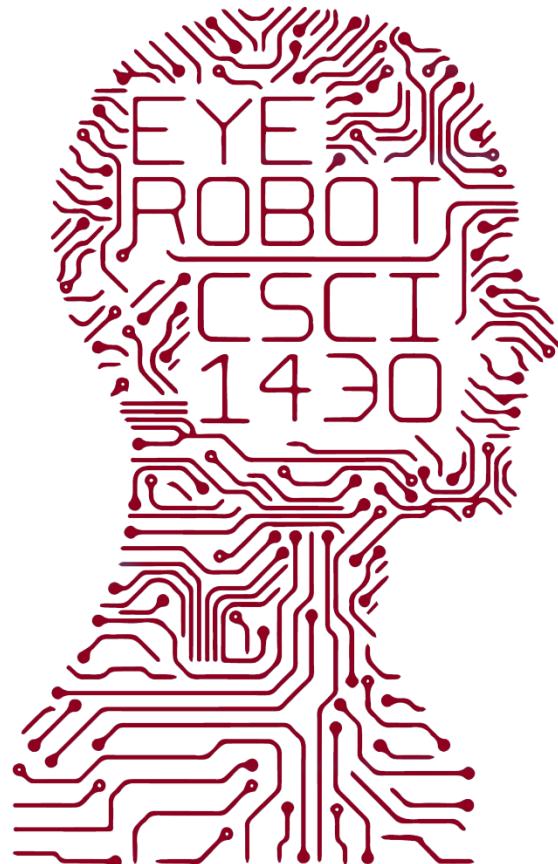


FUTURE VISION



COMPUTER VISION

Prismia.Chat <https://tinyurl.com/y26ragof>

Types of image transformations

What is an image?

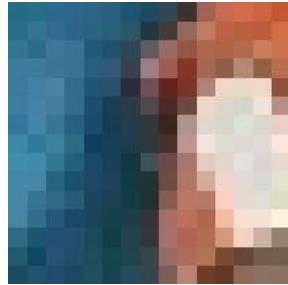


What is an image?



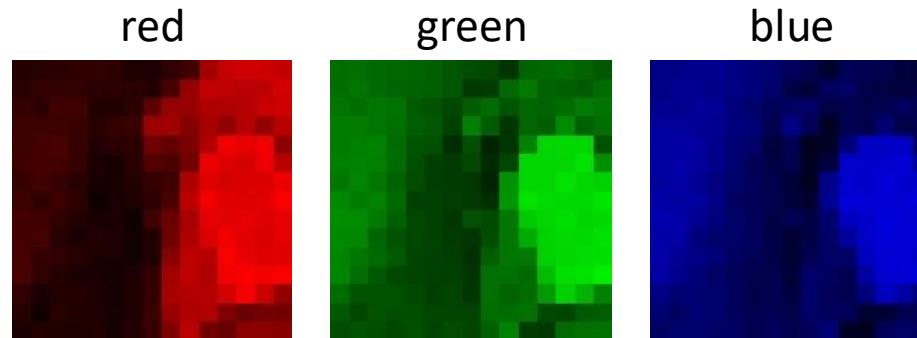
A (color) image
is a 3D tensor
of numbers.

What is an image?

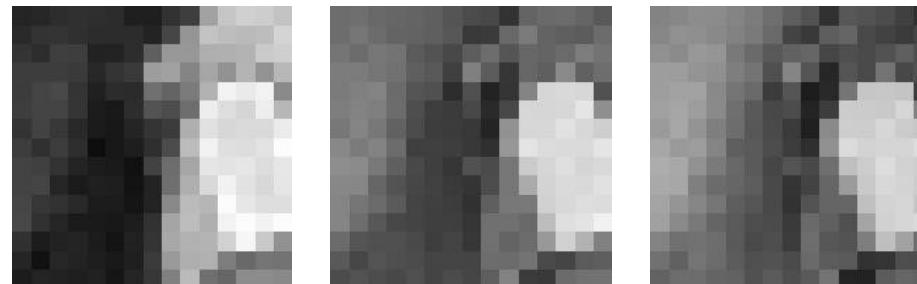


color image patch

How many bits are
the intensity values?



colorized for visualization



actual intensity values per channel

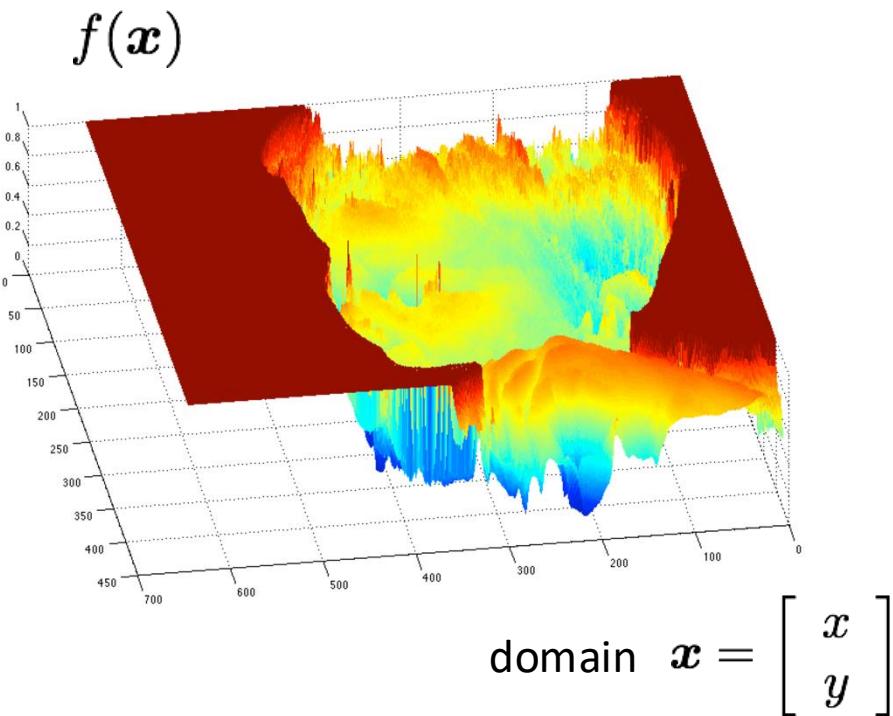
Each channel
is a 2D array of
numbers.

What is an image?



grayscale image

What is the range of
the image function f ?

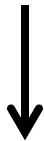


A (grayscale)
image is a 2D
function.

What types of image transformations can we do?



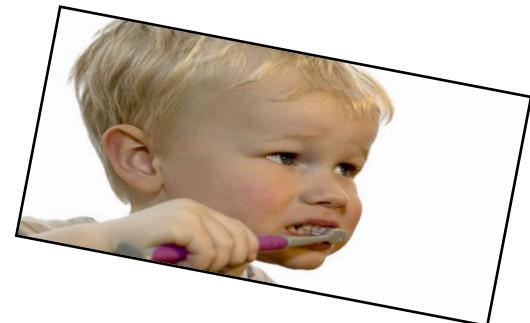
Filtering



changes pixel *values*



Warping



changes pixel *locations*

What types of image transformations can we do?

F



Filtering

$$G(\mathbf{x}) = h\{F(\mathbf{x})\}$$

G



changes *range* of image function

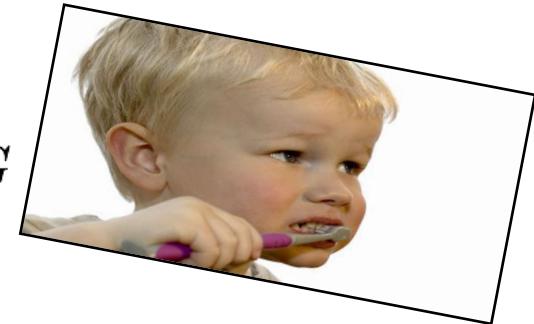
F



Warping

$$G(\mathbf{x}) = F(h\{\mathbf{x}\})$$

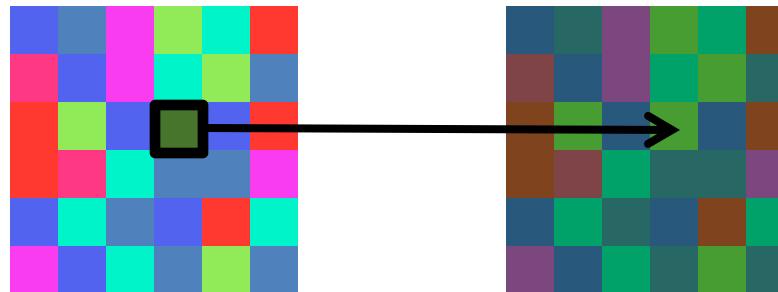
G



changes *domain* of image function

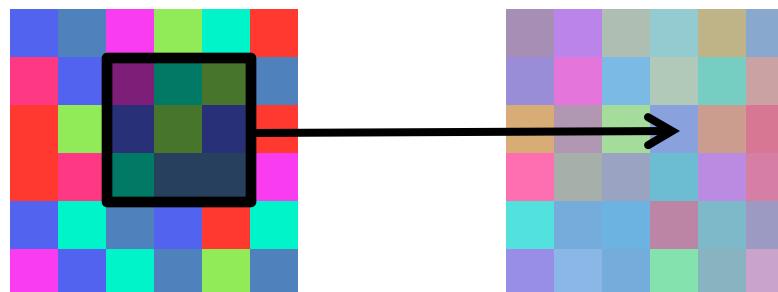
What types of image filtering can we do?

Point Operation



point processing

Neighborhood Operation



“filtering”

Point processing

Examples of point processing

original



darker



lower contrast



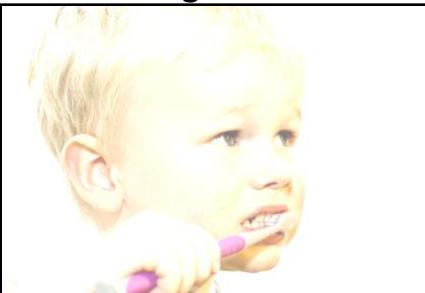
non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast



How would you
implement these?

Examples of point processing

original



darker



lower contrast



non-linear lower contrast

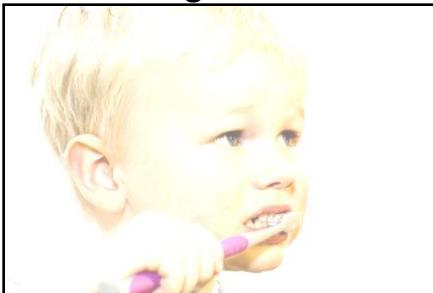


x

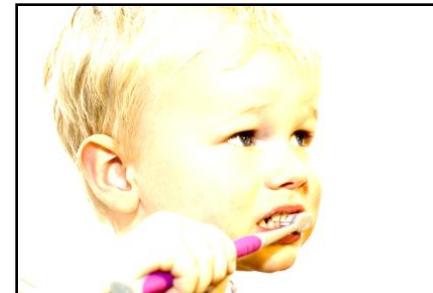
invert



lighten



raise contrast



non-linear raise contrast



How would you
implement these?

Examples of point processing

original



x

darker



$x - 128$

lower contrast



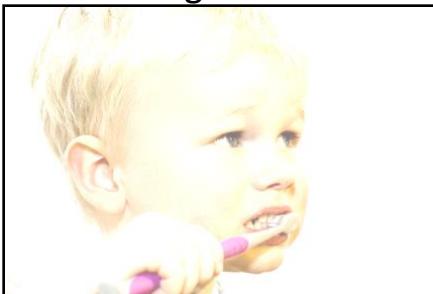
non-linear lower contrast



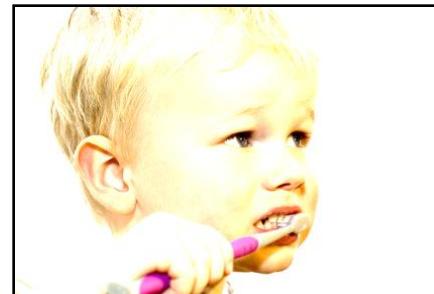
invert



lighten



raise contrast



non-linear raise contrast



How would you
implement these?

Examples of point processing

original



x

darker



$x - 128$

lower contrast



$\frac{x}{2}$

non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast



How would you
implement these?

Examples of point processing

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast

$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



lighten



raise contrast



non-linear raise contrast



How would you
implement these?

Examples of point processing

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast

$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



raise contrast



non-linear raise contrast



How would you
implement these?

Examples of point processing

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast

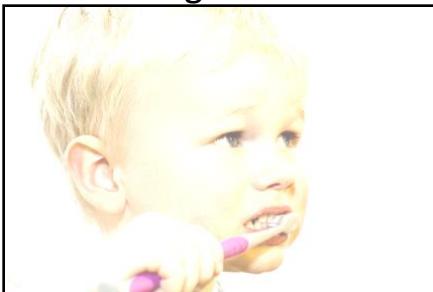
$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



non-linear raise contrast



How would you
implement these?

Examples of point processing

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast

$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



$$x \times 2$$

non-linear raise contrast



How would you
implement these?

Examples of point processing

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast

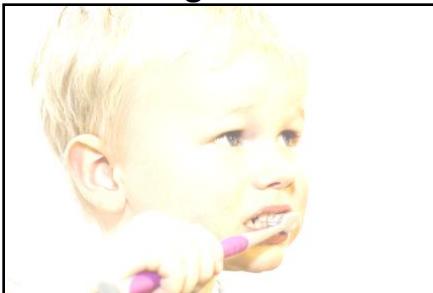
$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



$$x \times 2$$

non-linear raise contrast

$$\left(\frac{x}{255}\right)^2 \times 255$$

Many other types of point processing



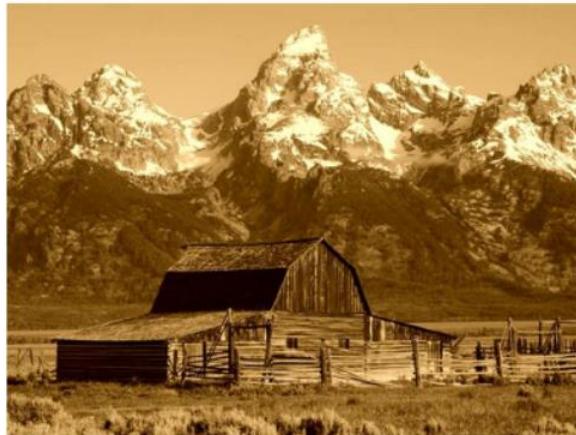
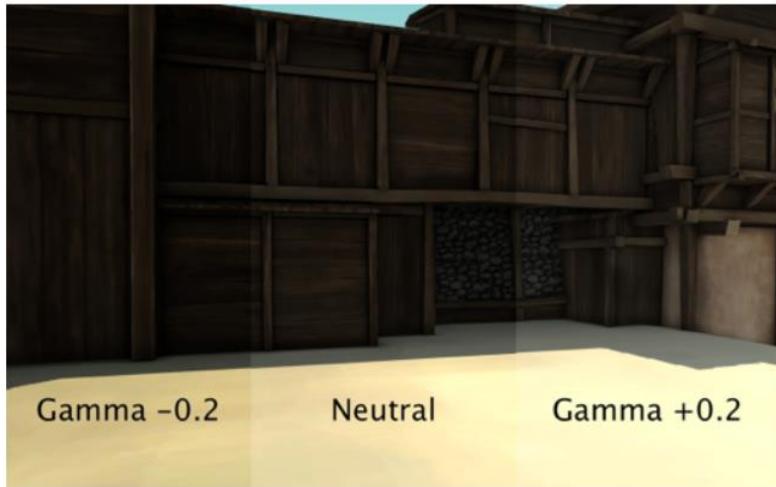
camera output



image after stylistic tonemapping

[Bae et al., SIGGRAPH 2006]

Many other types of point processing



Dimensionality of an Image

@ 8bit = 256 values \wedge 65,536

- Computer says ‘Inf’ combinations.

Some depiction of all possible scenes
would fit into this memory.

Dimensionality of an Image

@ 8bit = 256 values \wedge 65,536

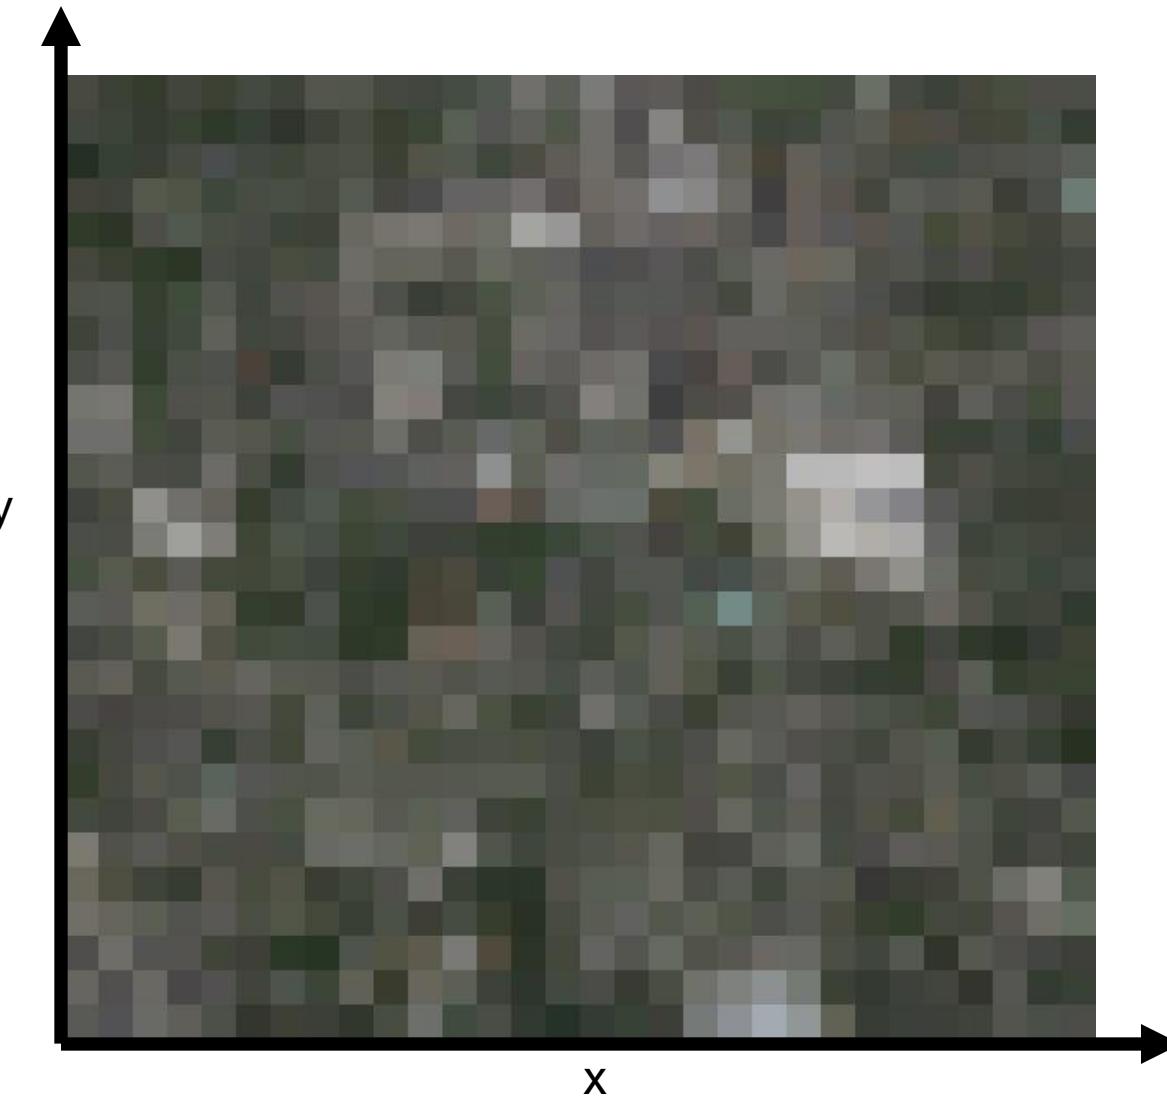
- Computer says ‘Inf’ combinations.

Some depiction of all possible scenes would fit into this memory.

Computer vision as making sense of an extremely high-dimensional space.

- Subspace of ‘natural’ images.
- Deriving low-dimensional, explainable models.

What is each part of an image?



What is each part of an image?

- Pixel -> picture element

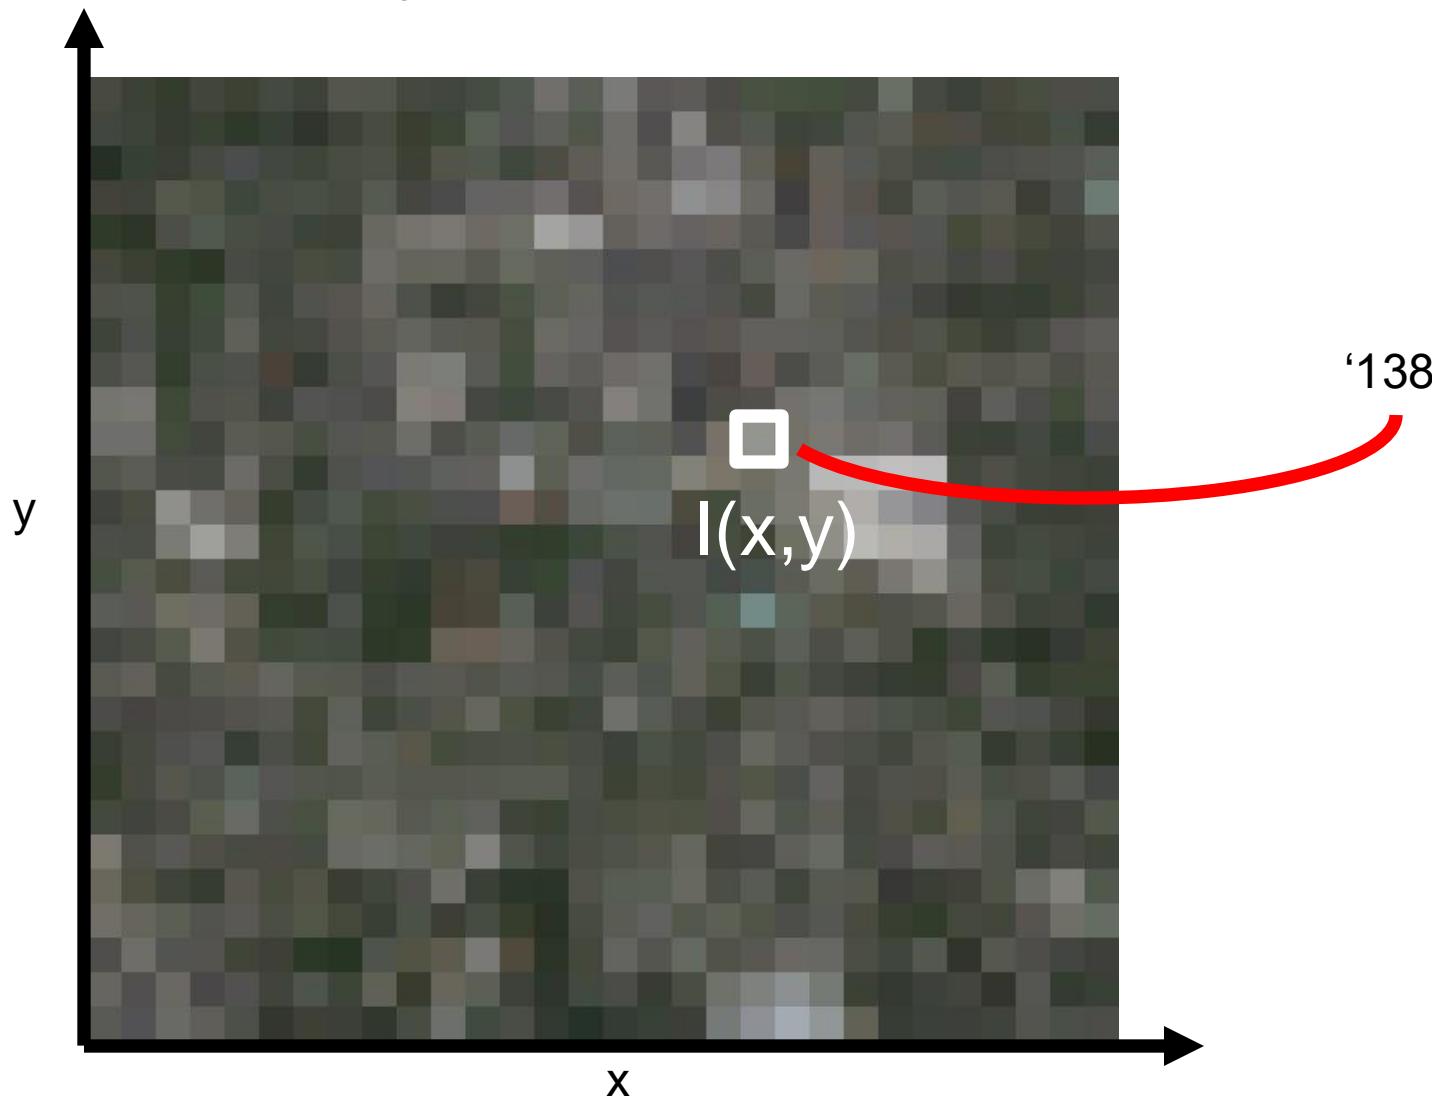
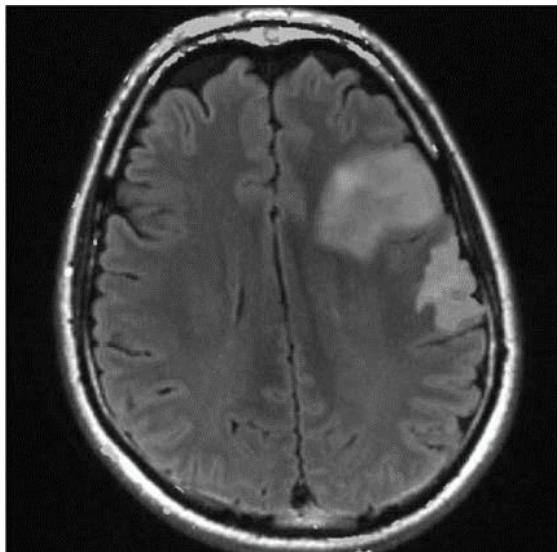


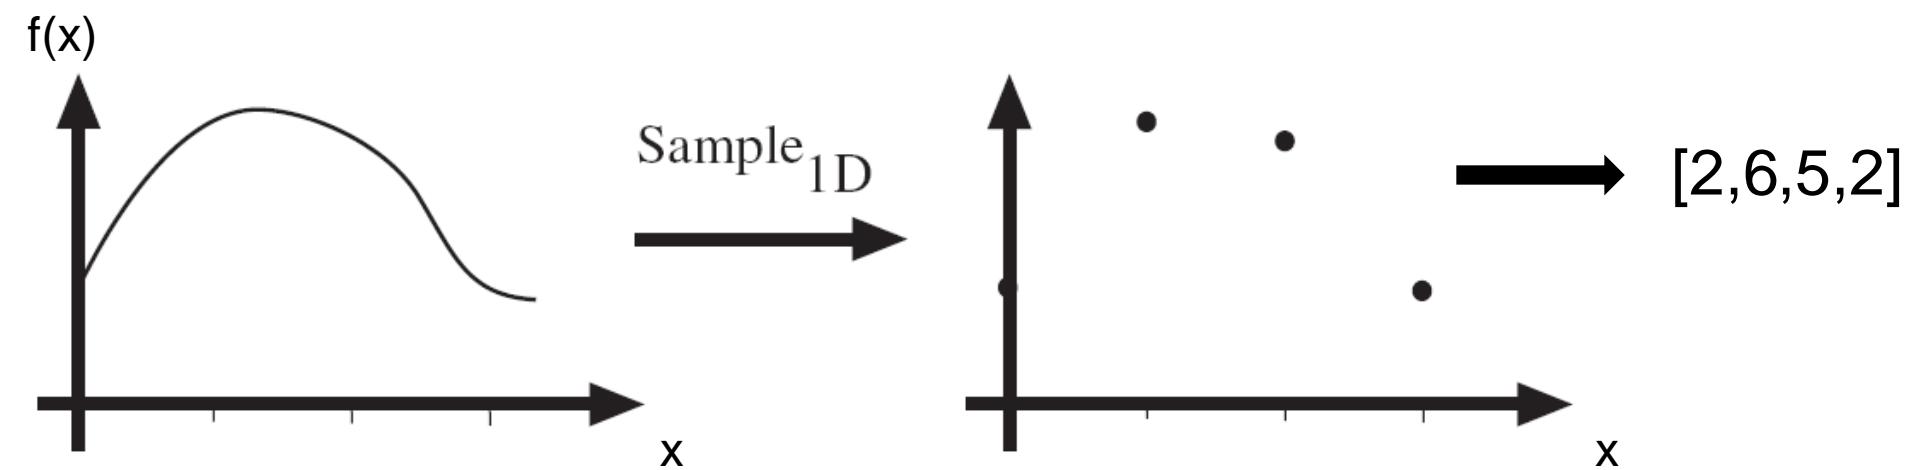
Image as a 2D sampling of signal

- Signal: function depending on some variable with physical meaning.
- Image: sampling of that function.
 - 2 variables: xy coordinates
 - 3 variables: xy + time (video)
 - ‘Brightness’ is the value of the function for visible light
- Can be other physical values too: temperature, pressure, depth ...

Example 2D Images

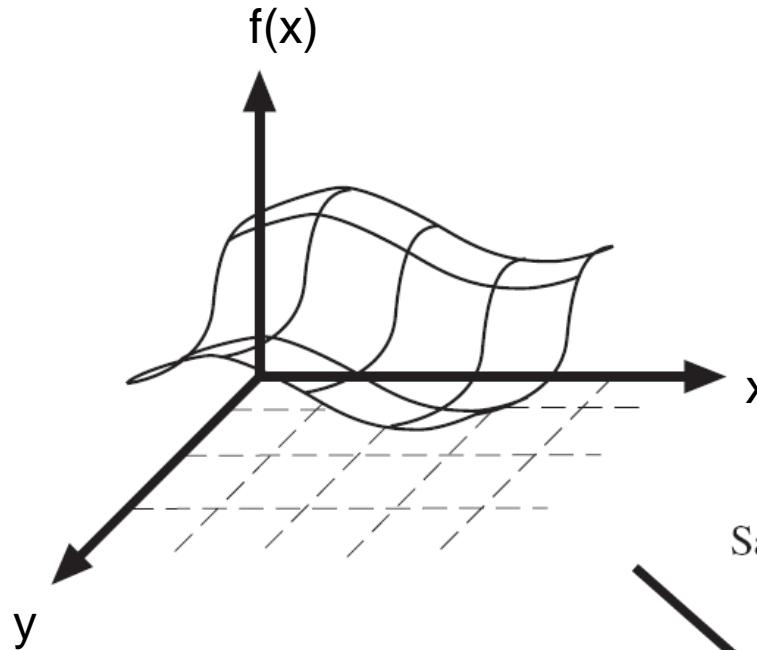


Sampling in 1D



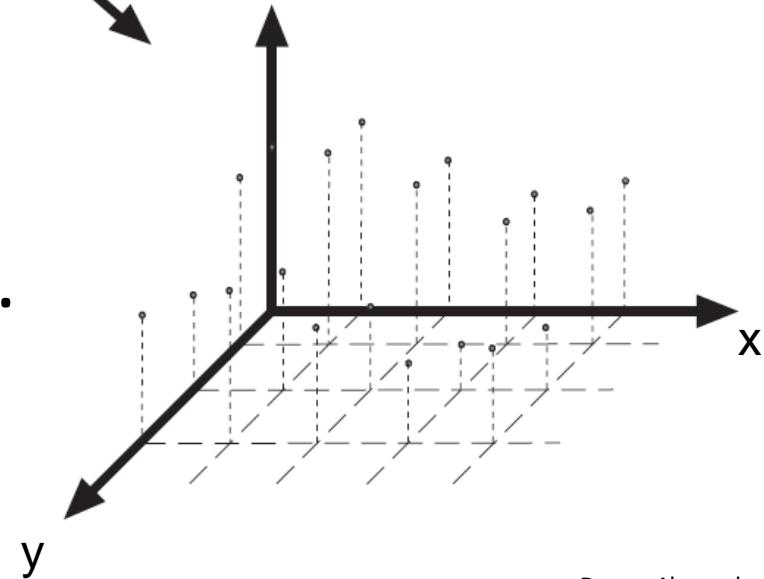
Sampling in 1D takes a function and returns a vector whose elements are values of that function at the sample points.

Sampling in 2D

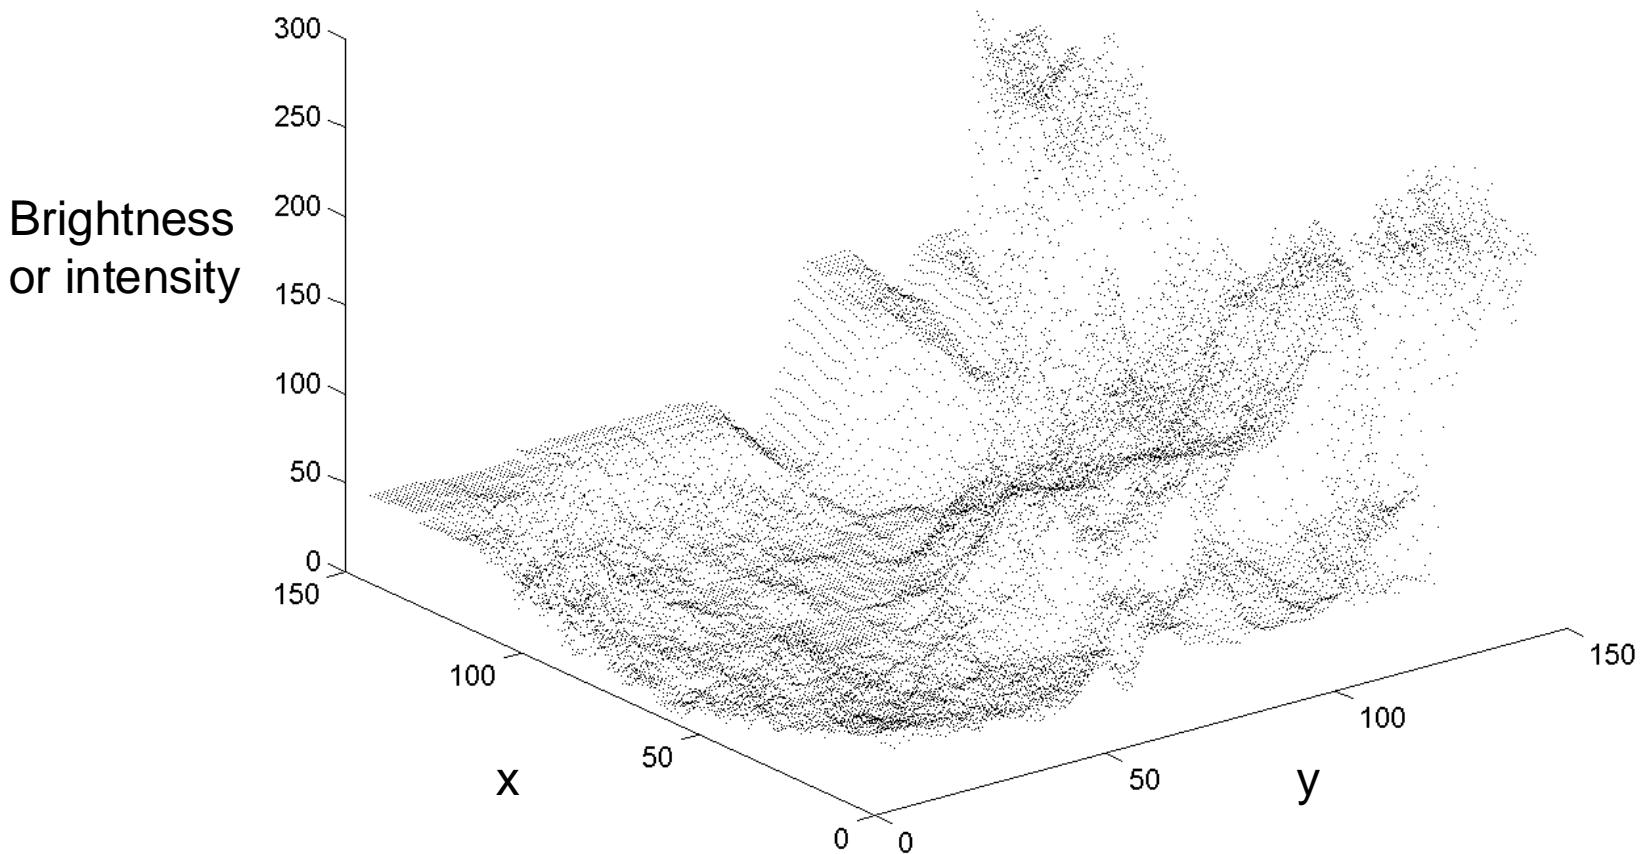


Sample_{2D}

Sampling in 2D takes a function and returns a matrix.

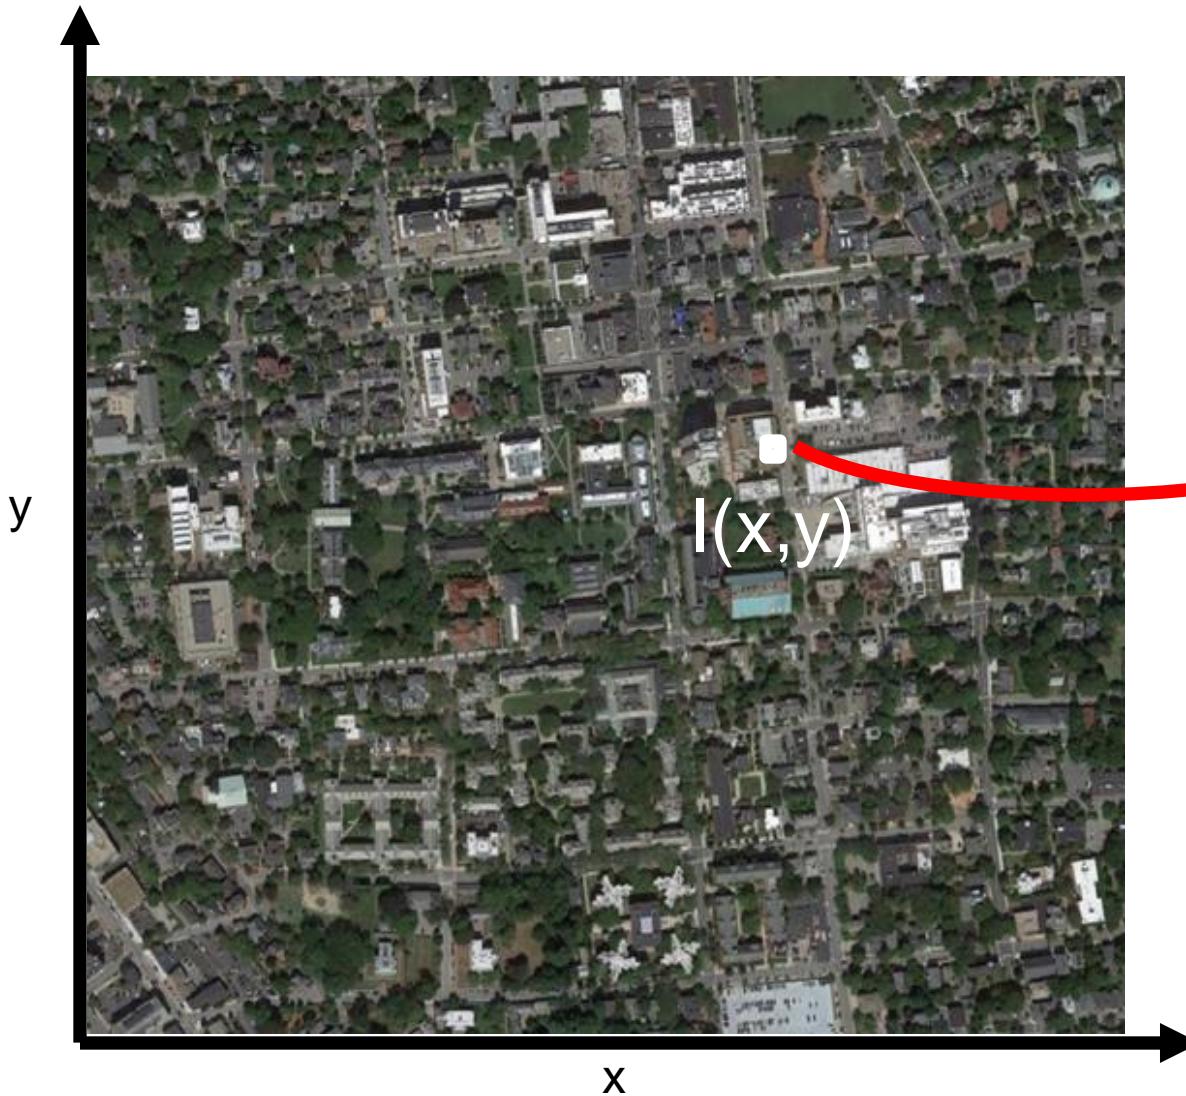


Grayscale Digital Image as ‘Height Map’

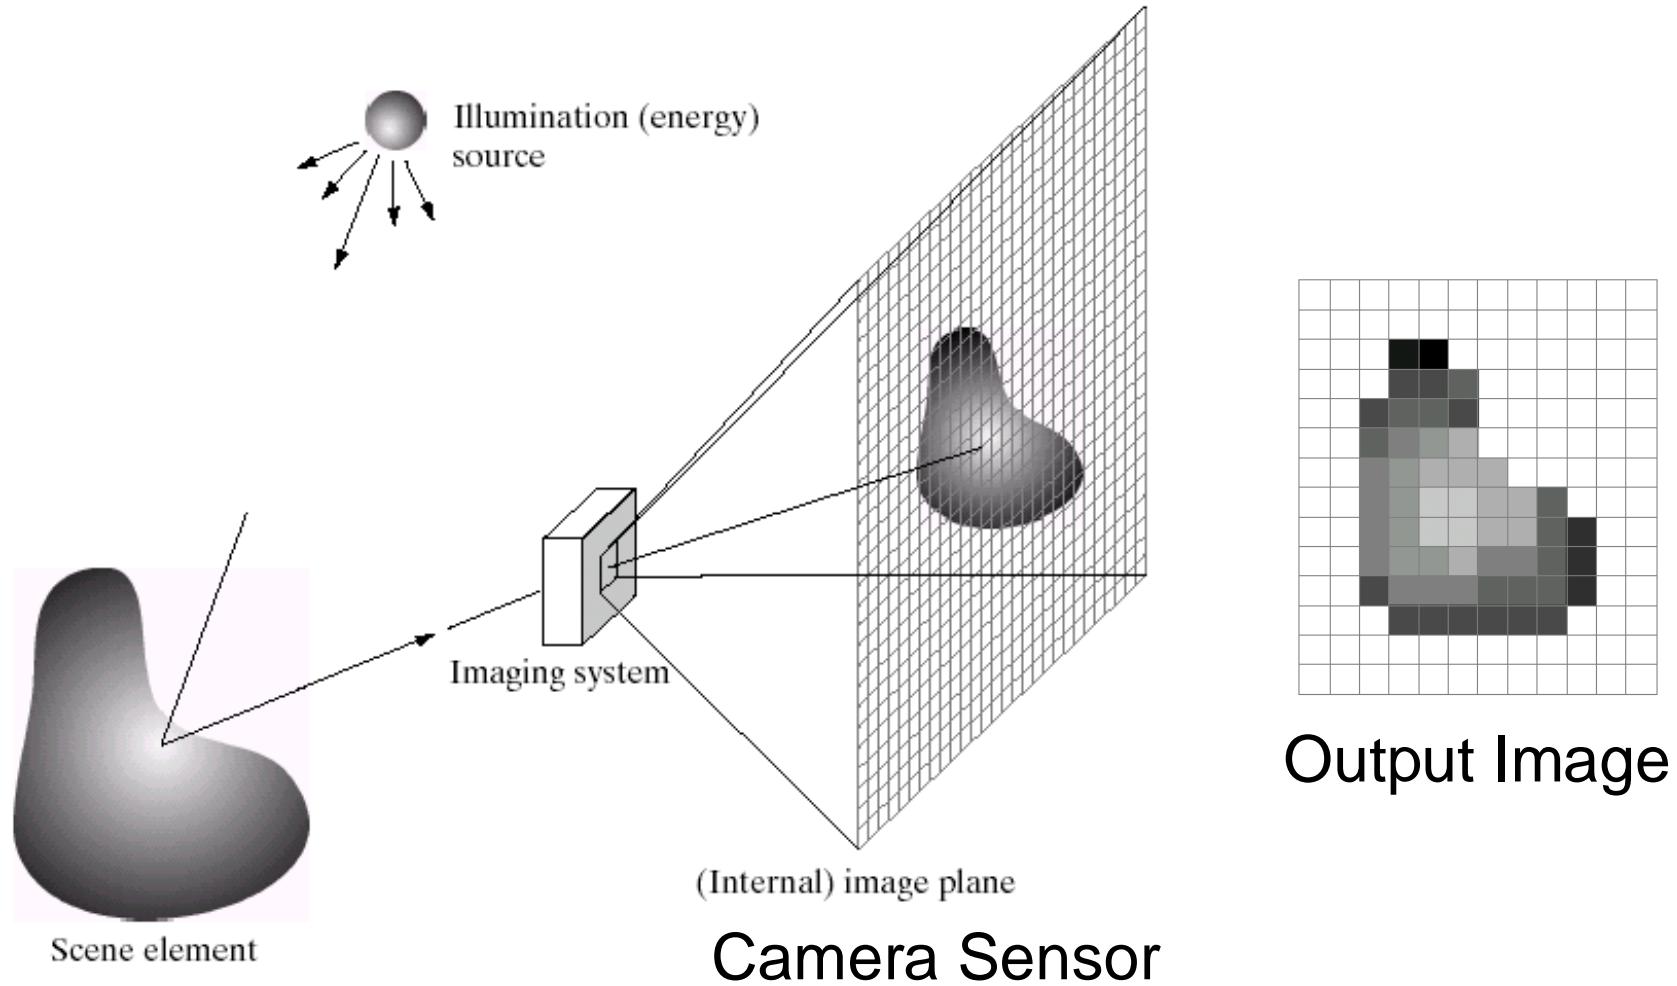


What is each part of a photograph?

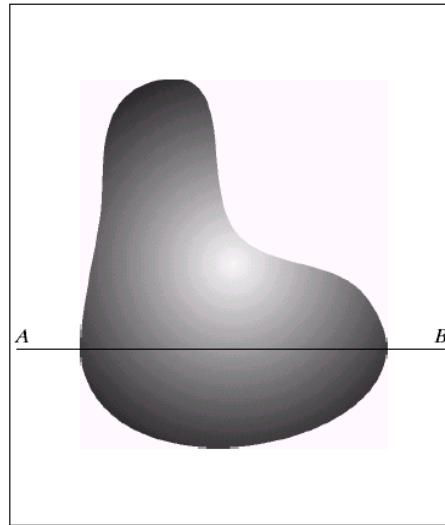
- Pixel -> picture element



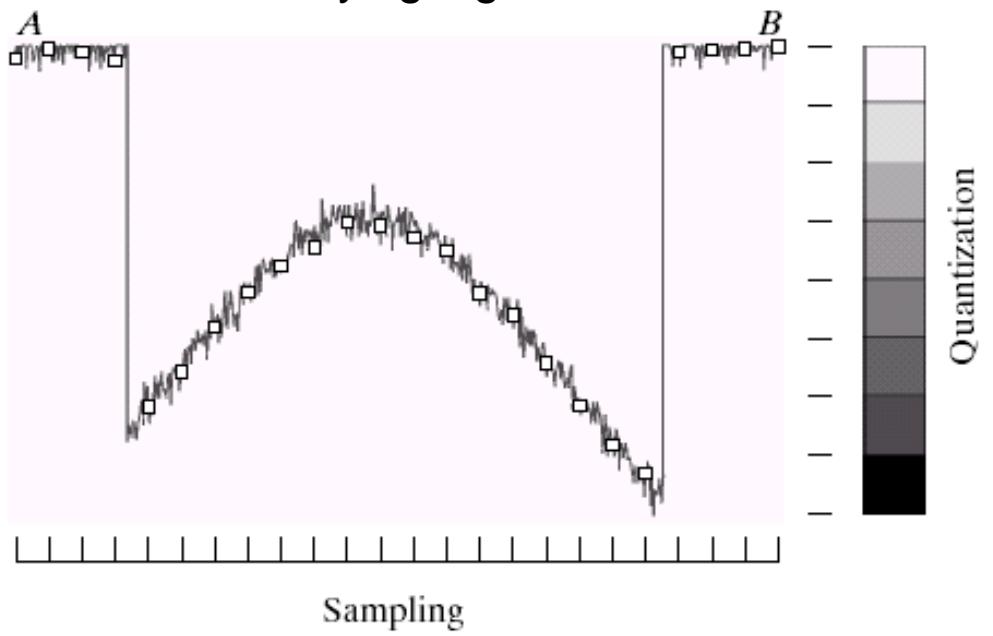
Integrating light over a range of angles.



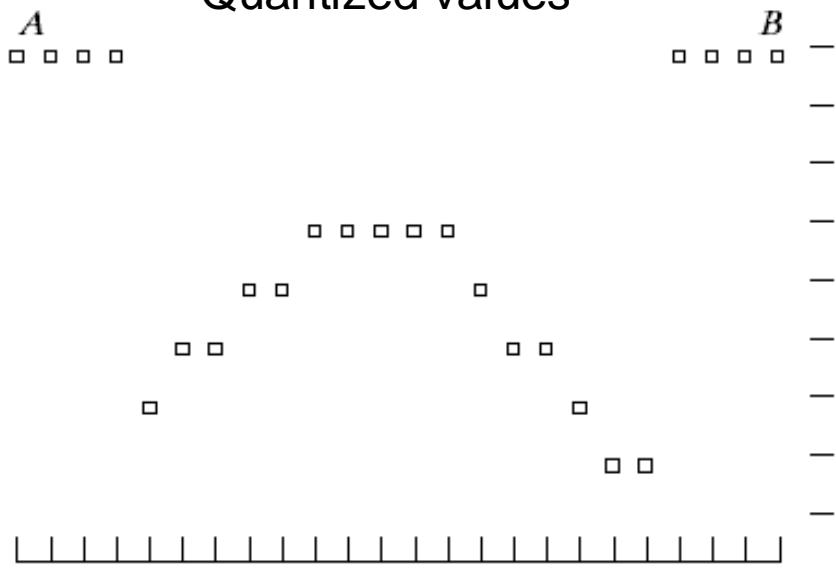
Quantization



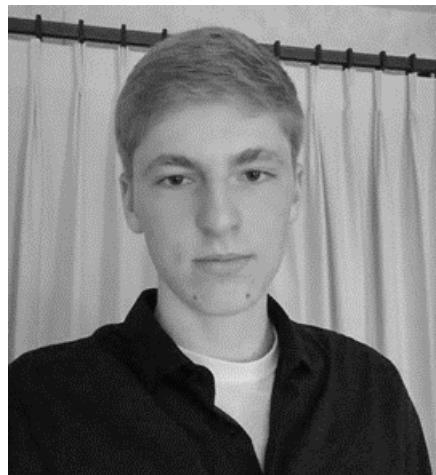
Underlying signal



Quantized values



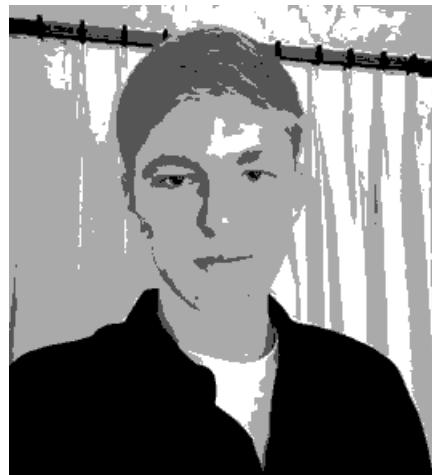
Quantization Effects – Radiometric Resolution



8 bit – 256 levels



4 bit – 16 levels



2 bit – 4 levels



1 bit – 2 levels

We often call this *bit depth*.
For photography, this is also related to *dynamic range*.

Linear shift-invariant image filtering

This week: three views of filtering

- Image filters in spatial domain
 - Filter is a mathematical operation of a grid of numbers
 - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
 - Filtering is a way to modify the frequencies of images
 - Denoising, sampling, image compression
- Image pyramids
 - Scale-space representation allows coarse-to-fine operations

Linear shift-invariant image filtering

- Replace each pixel by a *linear* combination of its neighbors (and possibly itself).
- The combination is determined by the filter's *kernel*.
- The same kernel is *shifted* to all pixel locations so that all pixels use the same linear combination of their neighbors.

Example: the box filter

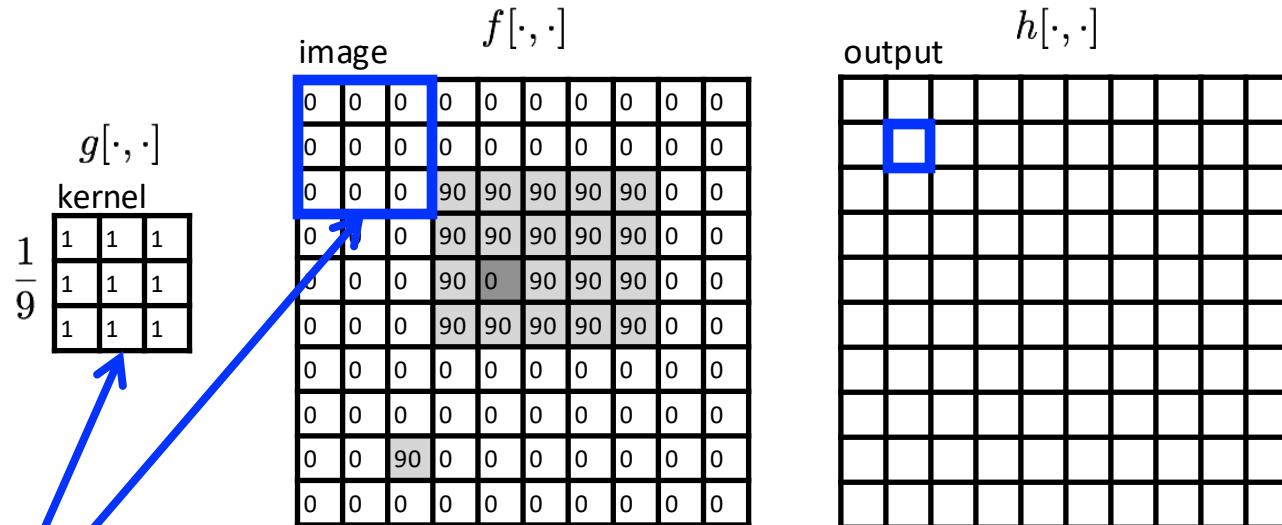
- also known as the 2D rect (not rekt) filter
- also known as the square mean filter

$$\text{kernel } g[\cdot, \cdot] = \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

- replaces pixel with local average
- has smoothing (blurring) effect



Let's run the box filter



note that we assume that
the kernel coordinates
are centered

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

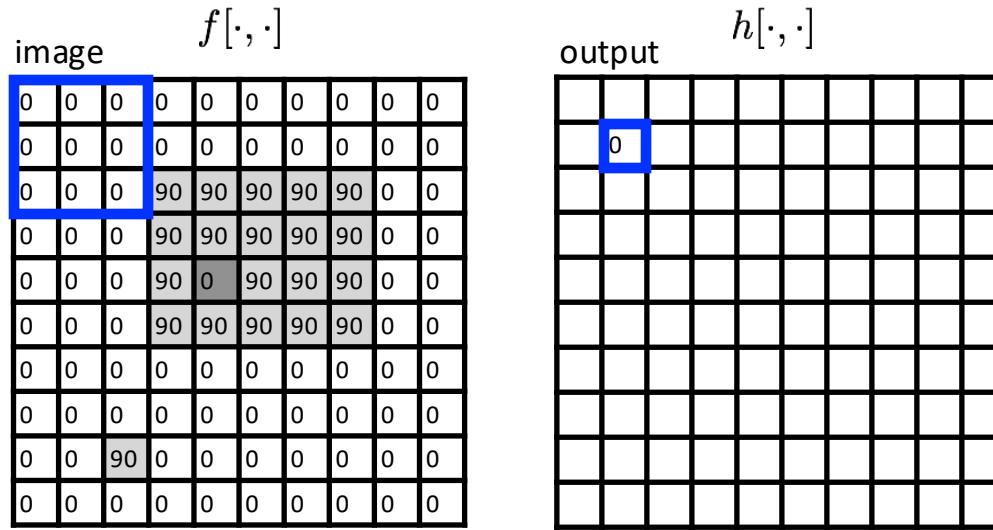
$$k, l = [-1, 0, 1]$$

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

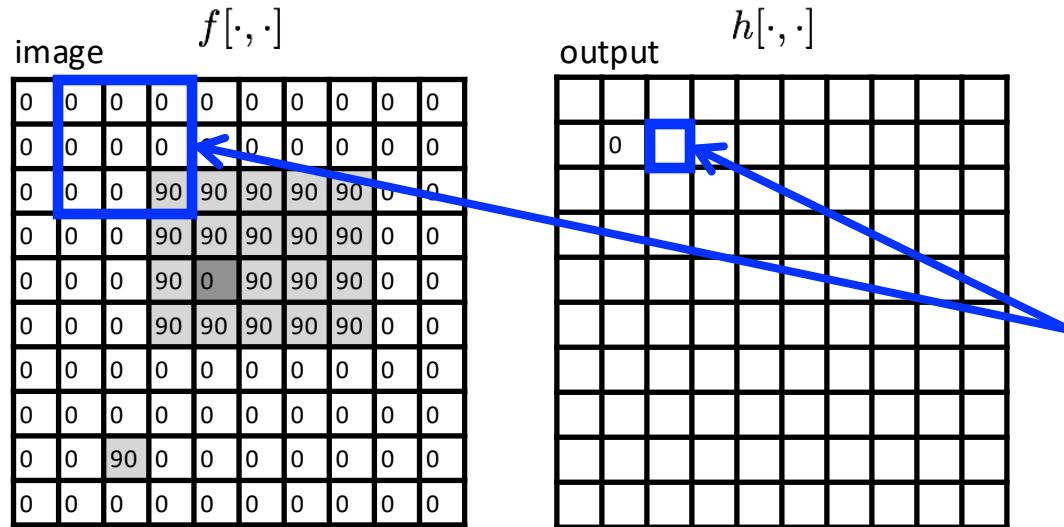
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



shift-invariant:
as the pixel
shifts, so does
the kernel

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

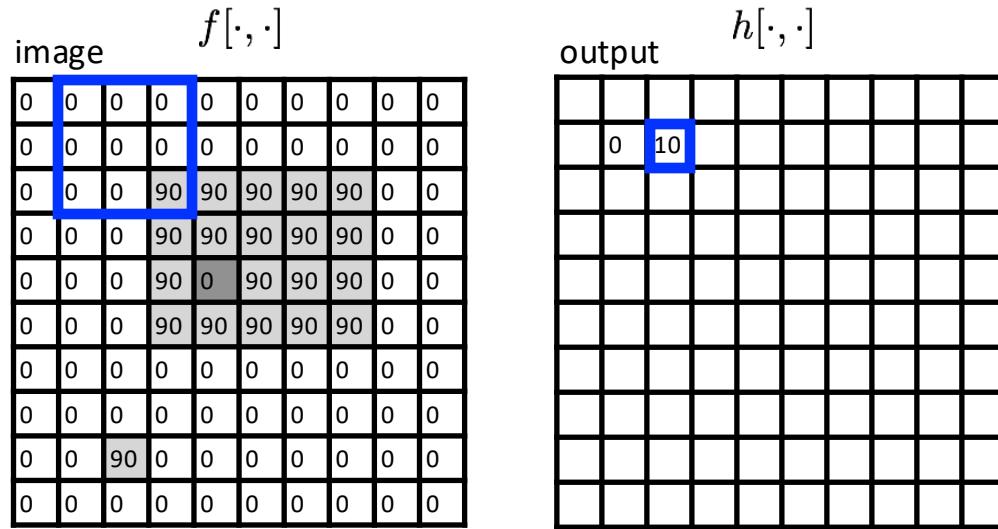
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

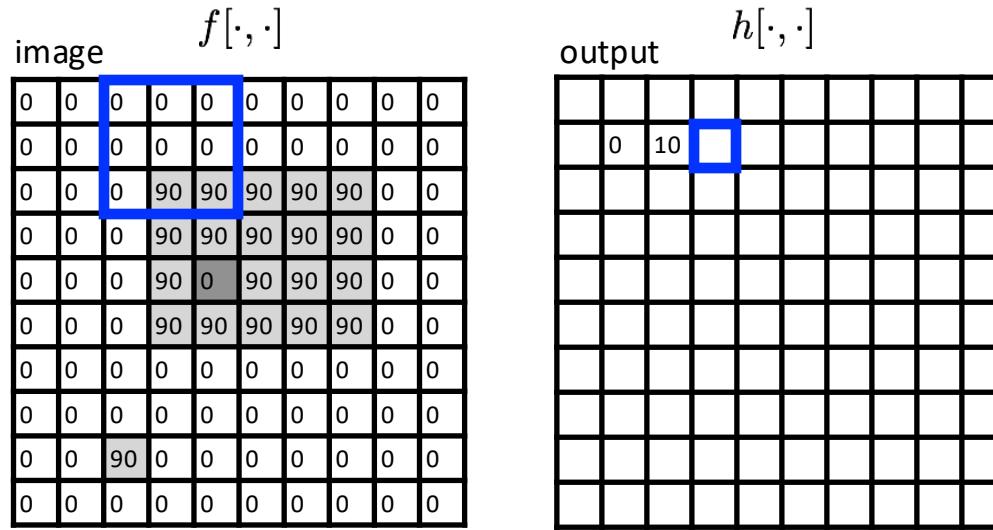
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

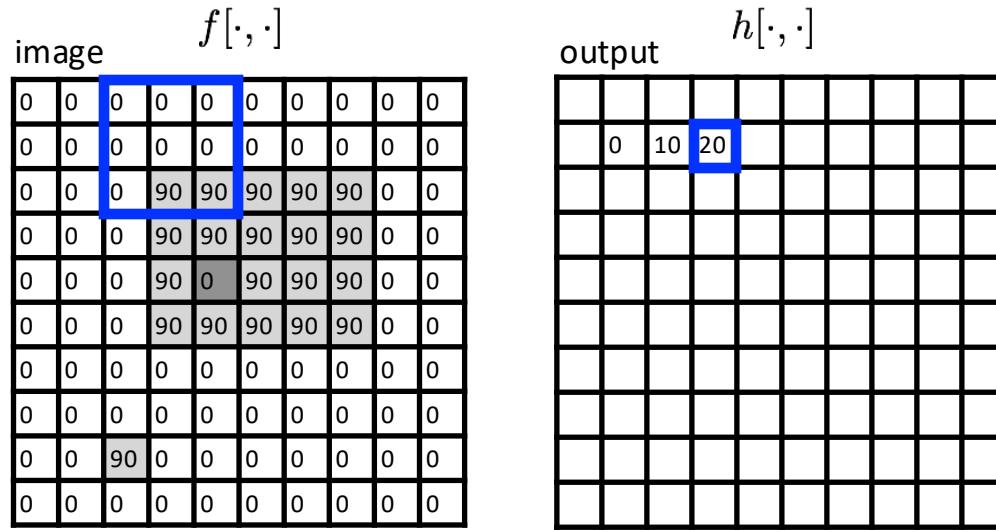
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

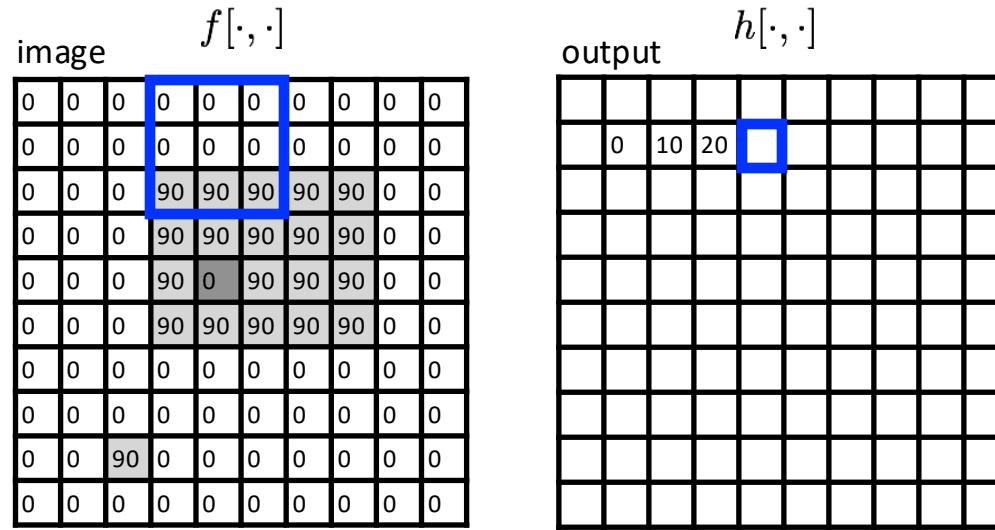
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

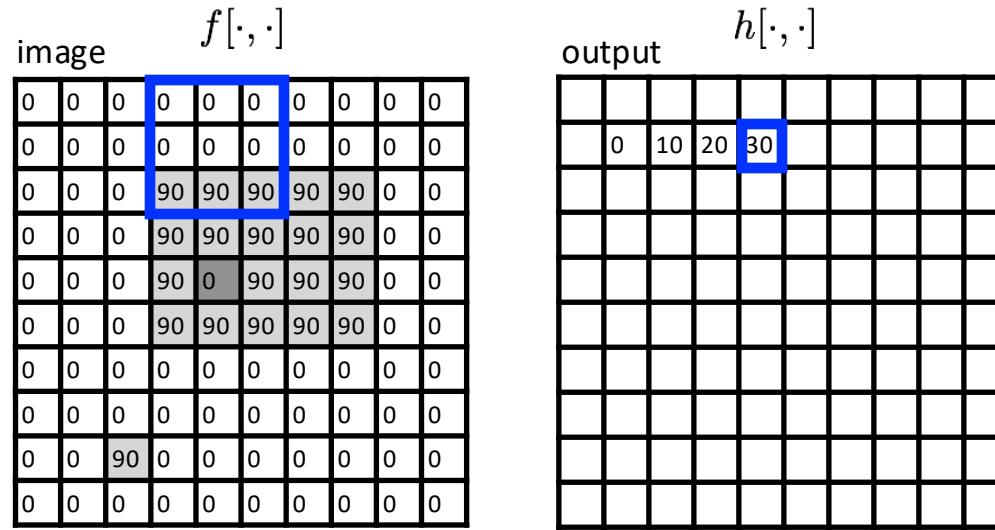
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

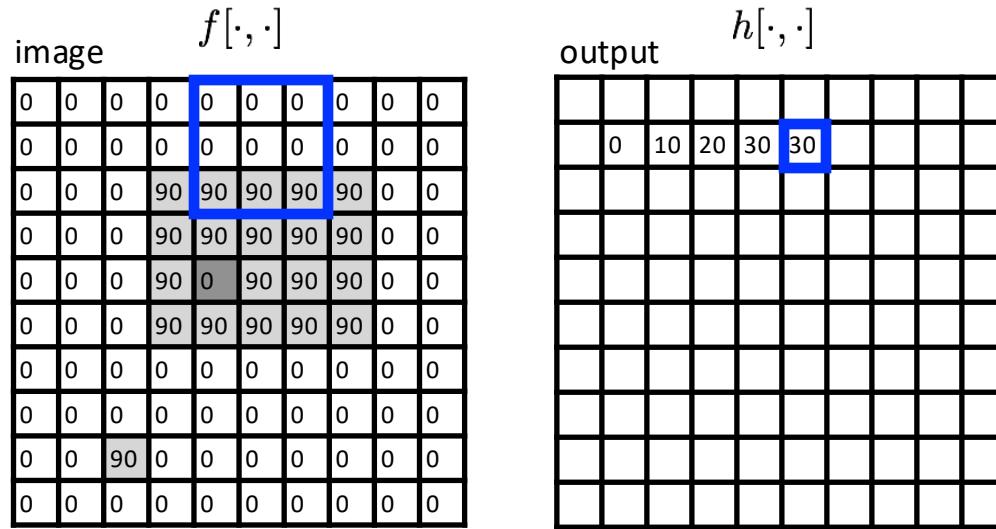
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

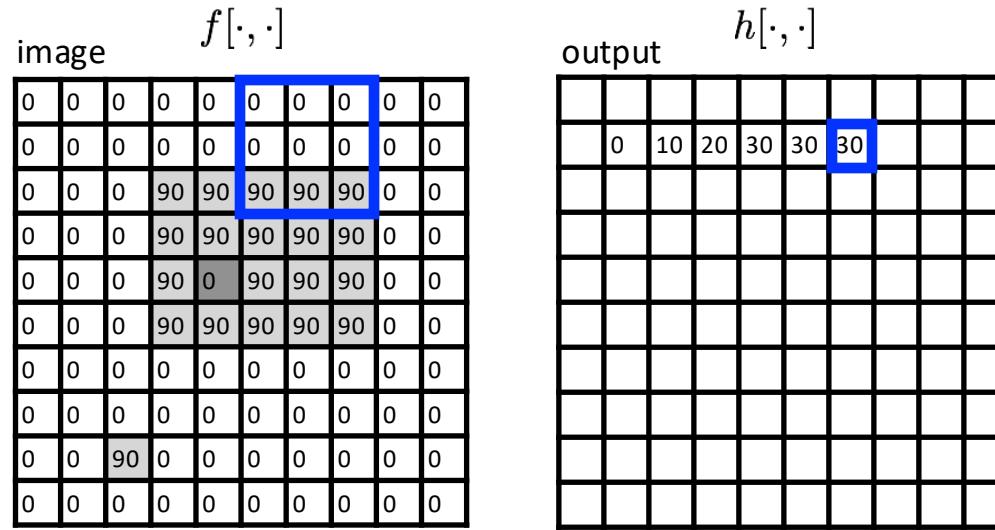
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

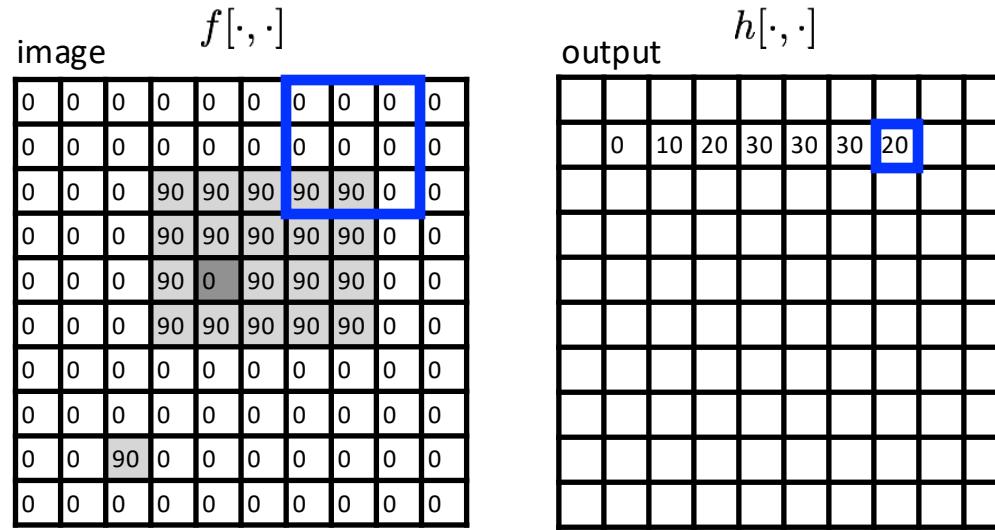
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

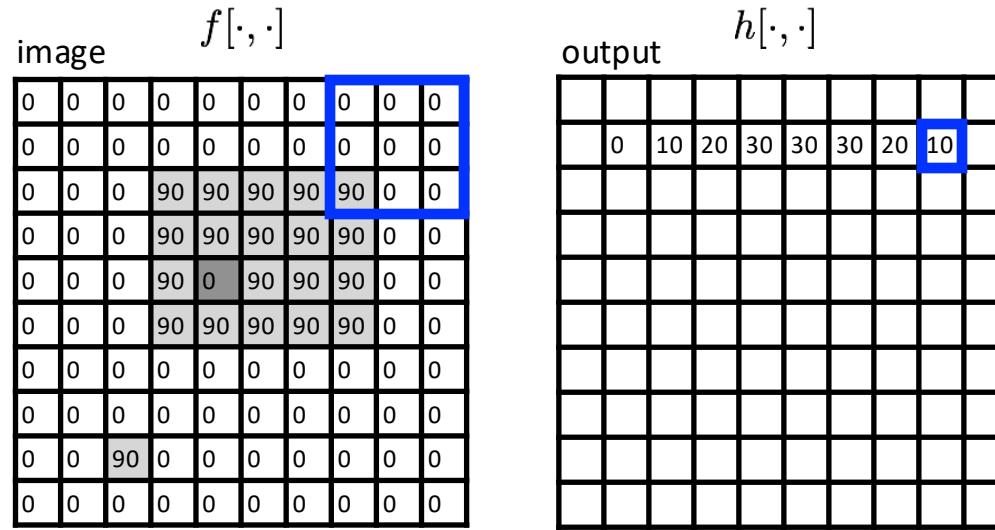
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

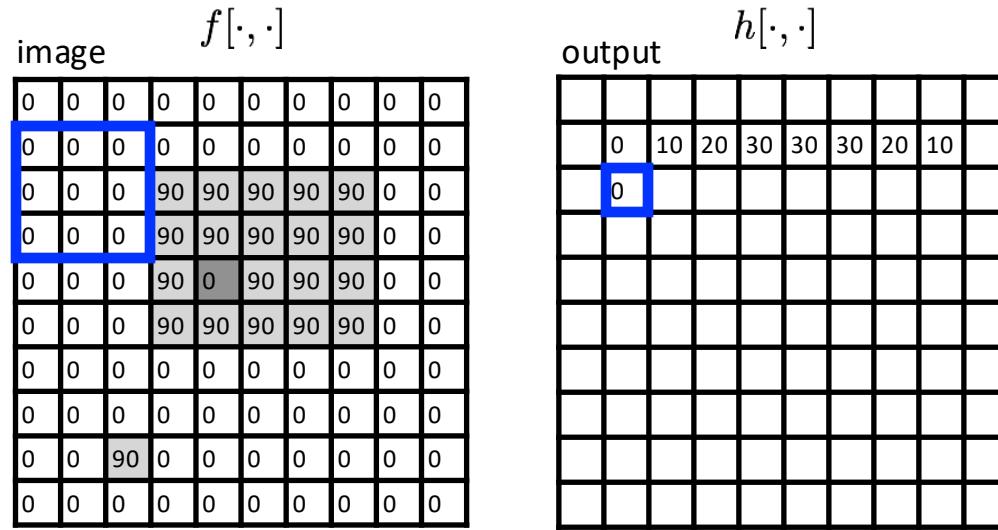
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output k, l filter image (signal)

Let's run the box filter

$$\begin{matrix} g[\cdot, \cdot] \\ \text{kernel} \\ \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \end{matrix}$$

A 10x10 grid representing an output matrix. The values are as follows:

	0	10	20	30	30	30	20	10	
0	0	10	20	30	30	30	20	10	
20	0	20	20	30	30	30	20	10	
	0	10	20	30	30	30	20	10	

The value at position (2, 3) is highlighted with a blue box.

$$h[m, n] = \sum_{k,l} g[k, l]f[m + k, n + l]$$

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

1	1	1
1	1	1
1	1	1

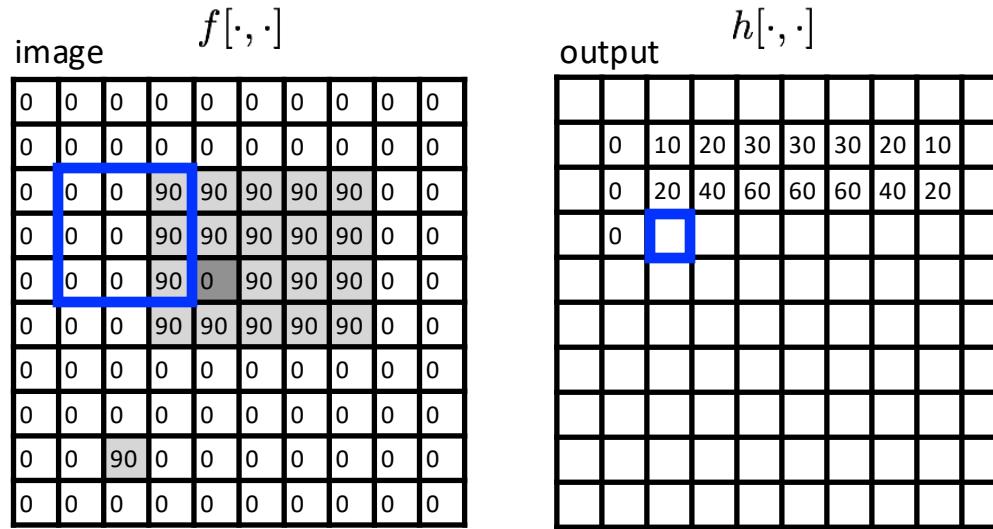
$$h[m, n] = \sum_{k,l} g[k, l]f[m + k, n + l]$$

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

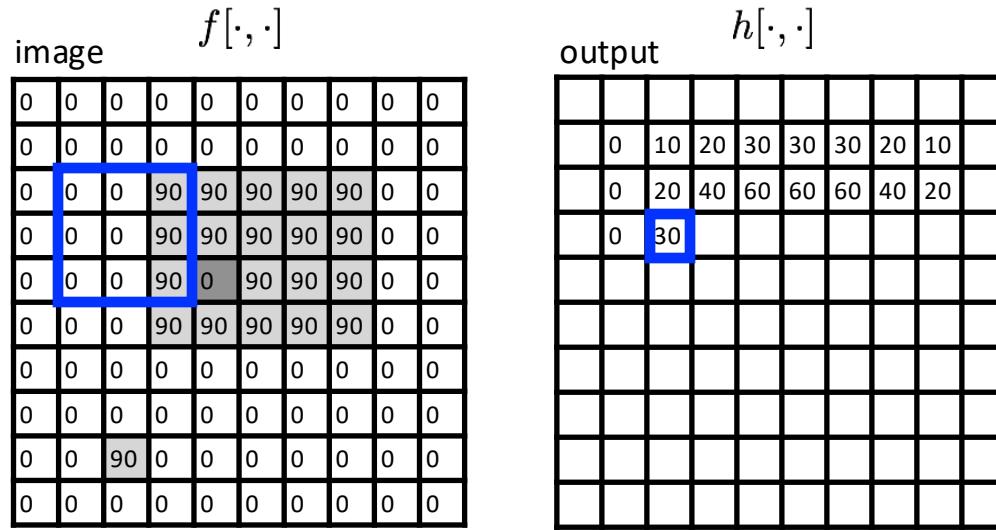
output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

image	$f[\cdot, \cdot]$	$h[\cdot, \cdot]$
0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10	
0 0 0 0 0 0 0 0 0 0	0 20 40 60 60 60 40 20	
0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30	
0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30	
0 0 0 90 0 90 90 90 0 0	0 20 30 50 50 60 40 20	
0 0 0 90 90 90 90 90 0 0	0 10 20 30 30 30 20 10	
0 0 0 0 0 0 0 0 0 0	10 10 10 10 0 0 0 0	
0 0 90 0 0 0 0 0 0 0	10	
0 0 0 0 0 0 0 0 0 0		

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output k, l filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

image	$f[\cdot, \cdot]$	$h[\cdot, \cdot]$
0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10	
0 0 0 0 0 0 0 0 0 0	0 20 40 60 60 60 40 20	
0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30	
0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30	
0 0 0 90 0 90 90 90 0 0	0 20 30 50 50 60 40 20	
0 0 0 90 90 90 90 90 0 0	0 10 20 30 30 30 20 10	
0 0 0 0 0 0 0 0 0 0	10 10 10 10 0 0 0 0	
0 0 90 0 0 0 0 0 0 0	10 10 10 10 0 0 0 0	0
0 0 0 0 0 0 0 0 0 0		

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output k, l filter image (signal)

... and the result is

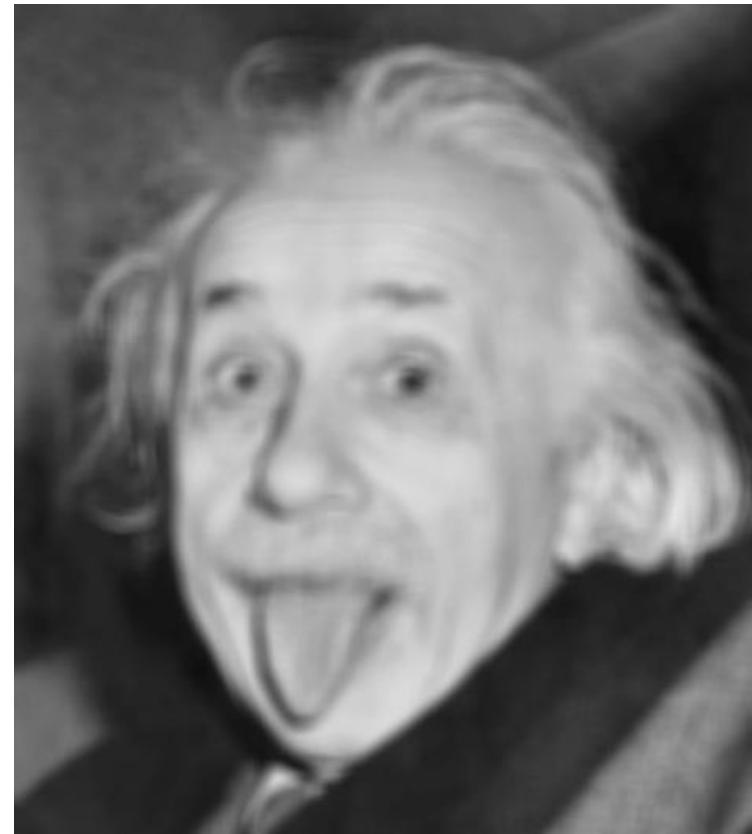
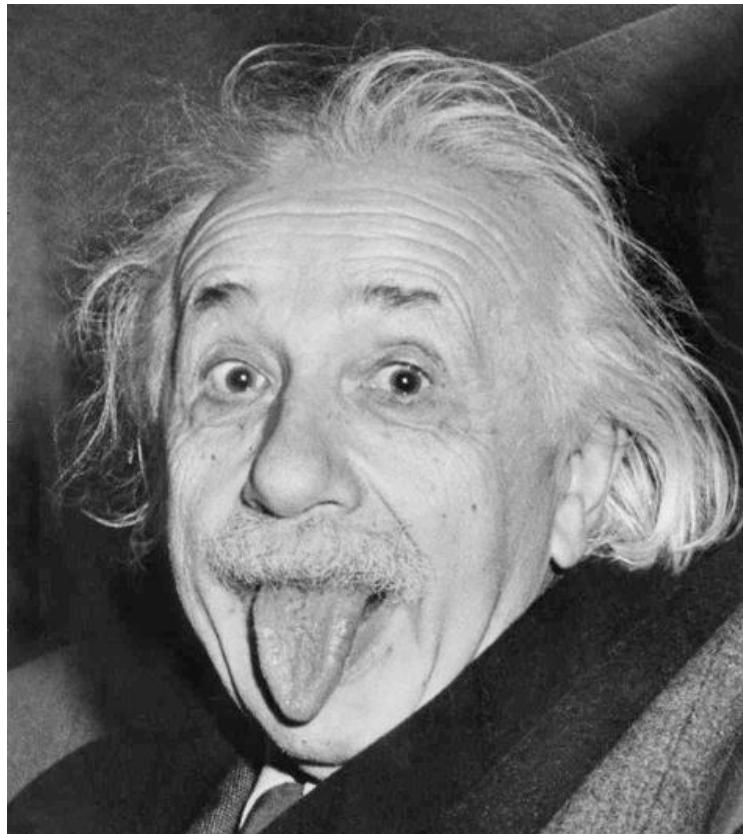
$$\frac{1}{9} \begin{matrix} g[\cdot, \cdot] \\ \text{kernel} \\ \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \end{matrix}$$

image	$f[\cdot, \cdot]$	$h[\cdot, \cdot]$
output		
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10
0 0 0 0 0 0 0 0 0 0	0 0 0 90 90 90 90 90 0 0	0 20 40 60 60 60 40 20
0 0 0 90 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30
0 0 0 90 0 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30
0 0 0 90 90 90 90 90 0 0	0 0 0 0 0 0 0 0 0 0	0 20 30 50 50 60 40 20
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10
0 0 0 0 0 0 0 0 0 0	0 0 90 0 0 0 0 0 0 0	10 10 10 10 0 0 0 0
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	10 10 10 10 0 0 0 0

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output k, l filter image (signal)

Some more realistic examples



Some more realistic examples



Some more realistic examples

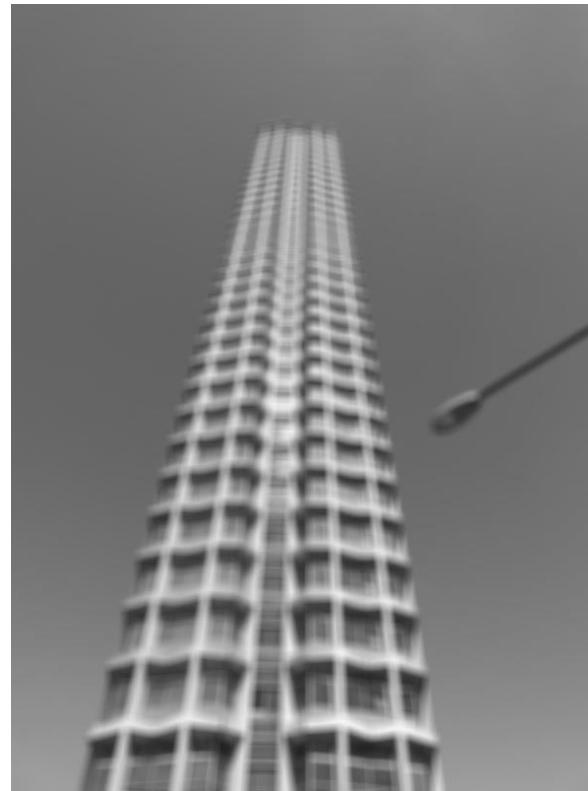


Image filtering

Compute function of local neighborhood
at each position:

$$h[m, n] = \sum_{k,l} f[k, l] I[m+k, n+l]$$

Image filtering

Compute function of local neighborhood
at each position:

h=output

f=filter

I=image

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

2d coords=k, l 2d coords=m, n

[]

[]

[]

Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove ‘sharp’ features)

$$f[\cdot, \cdot]$$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove ‘sharp’ features)
- Why does it sum to one?

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Smoothing with box filter

 $f[\cdot, \cdot]$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

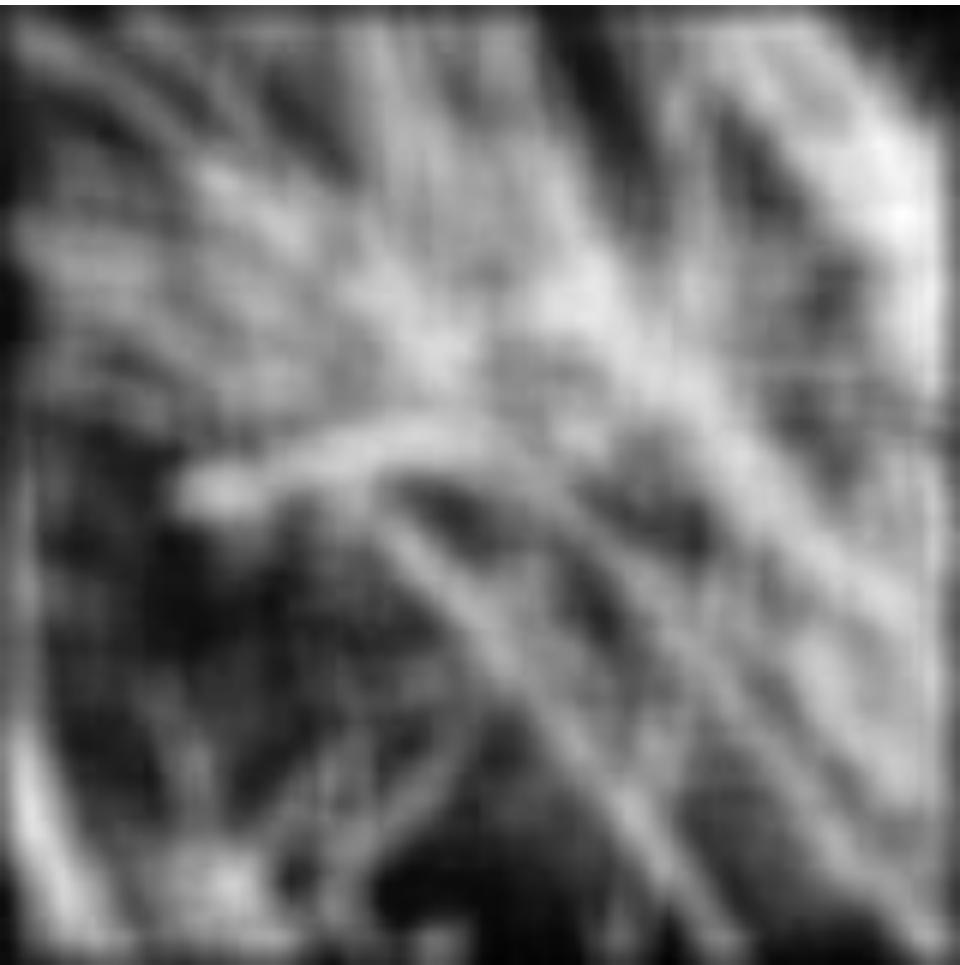


Image filtering

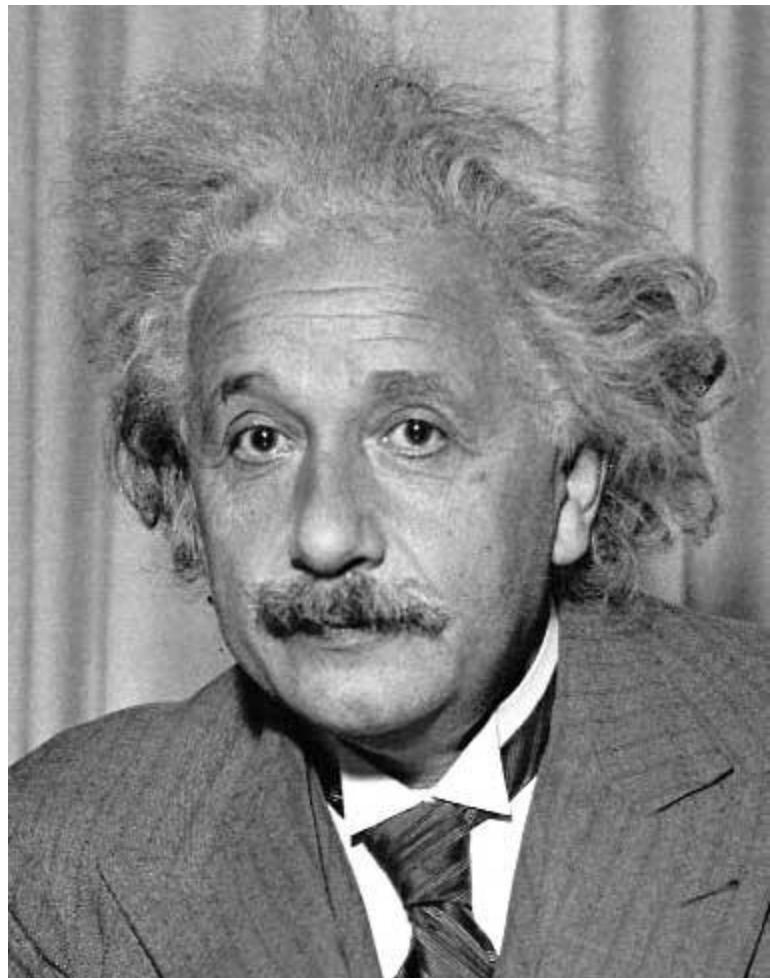
Compute function of local neighborhood
at each position:

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Really important!

- Enhance images
 - Denoise, resize, increase contrast, etc.
- Extract information from images
 - Texture, edges, distinctive points, etc.
- Detect patterns
 - Template matching

Filtering practice





- 1.
- | | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |
- 2.
- | | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 0 |
- 3.
- | | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |
- 4.
- | | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 2 | 0 |
| 0 | 0 | 0 |
- $\quad - \quad \frac{1}{9}$
- | | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

1. Practice with linear filters



0	0	0
0	1	0
0	0	0

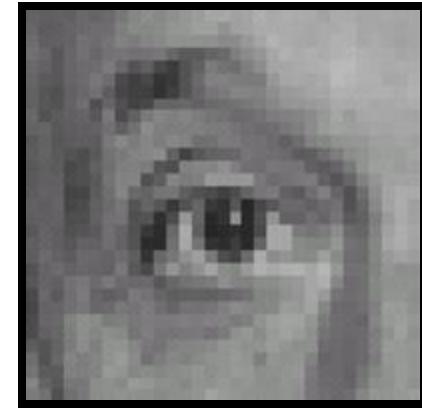
?

1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

2. Practice with linear filters



0	0	0
0	0	1
0	0	0

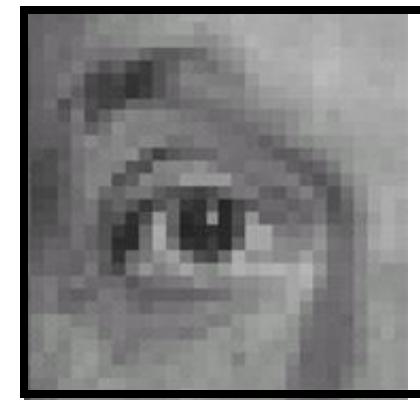
?

2. Practice with linear filters



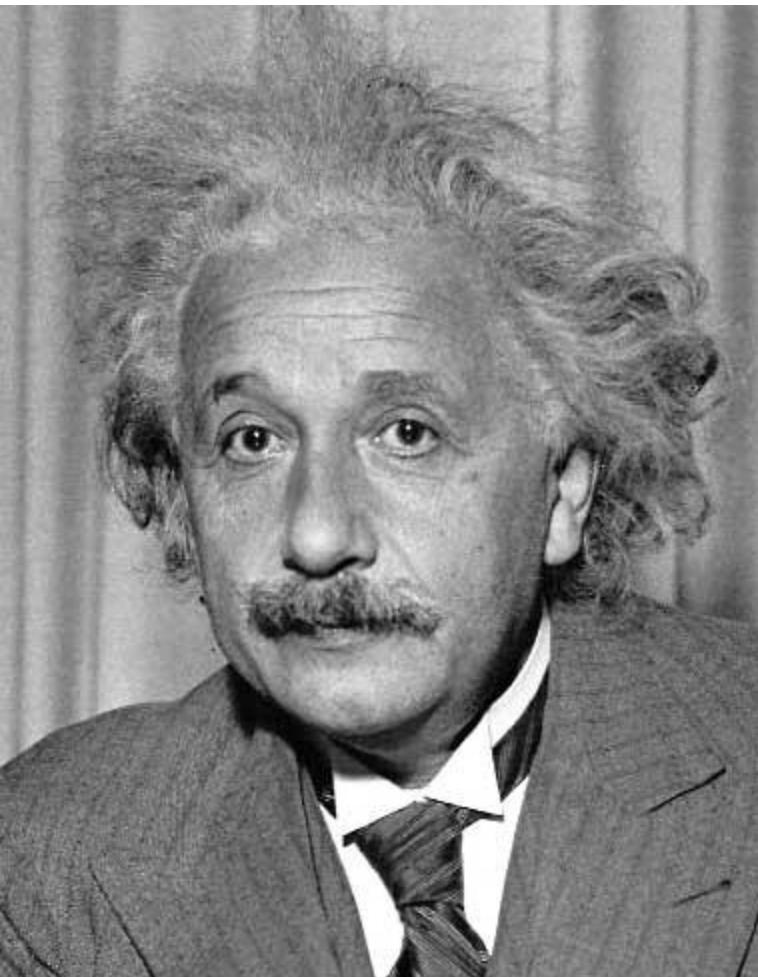
Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

3. Practice with linear filters



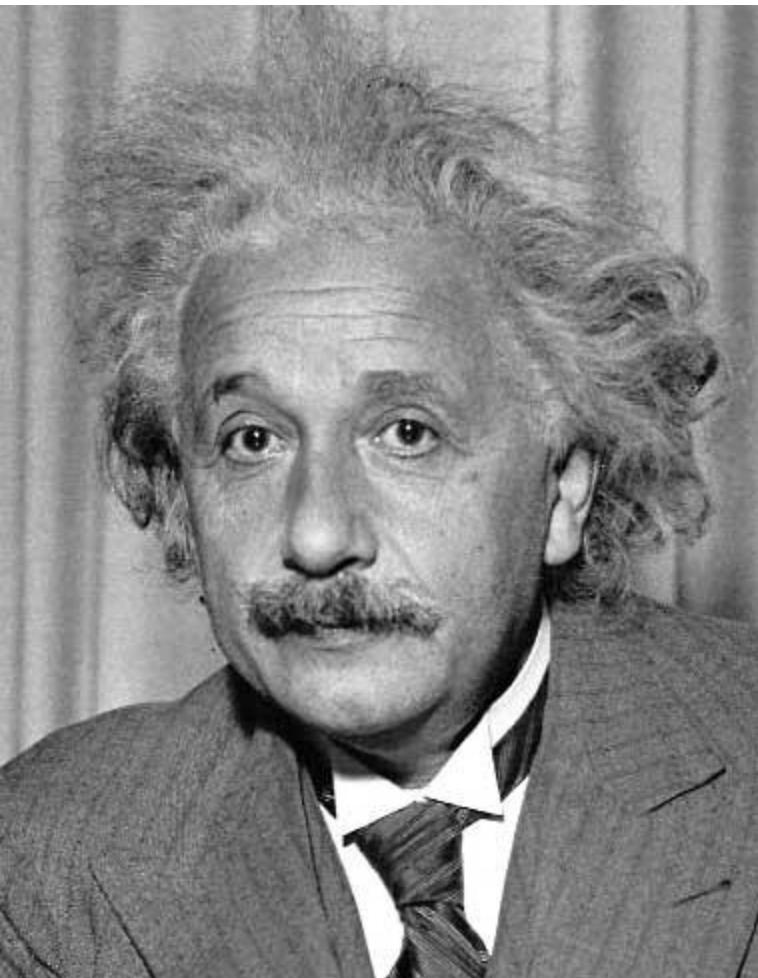
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

3. Practice with linear filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$$\frac{1}{9}$$

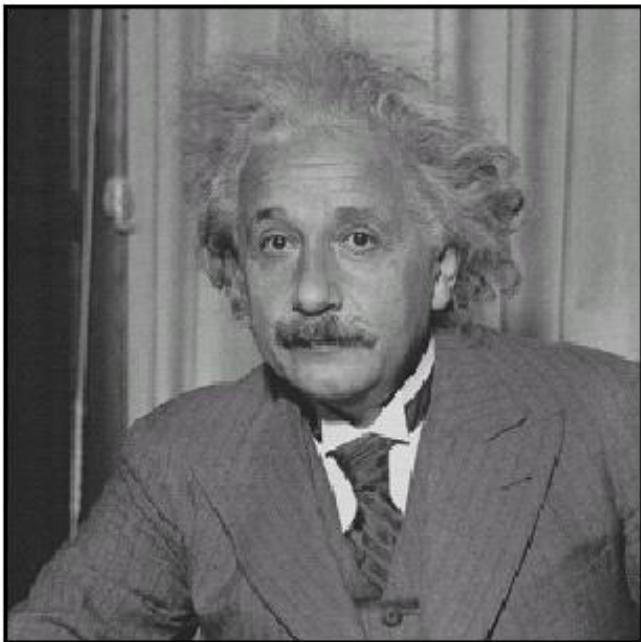
1	1	1
1	1	1
1	1	1



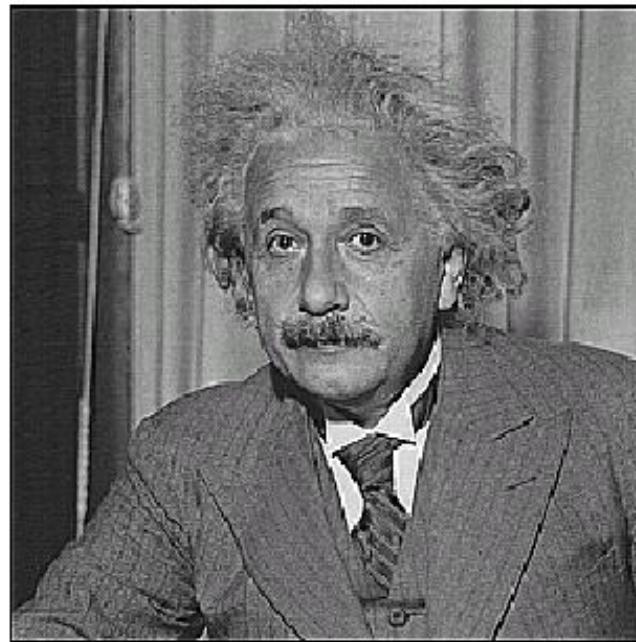
Sharpening filter

- Accentuates differences with local average

4. Practice with linear filters

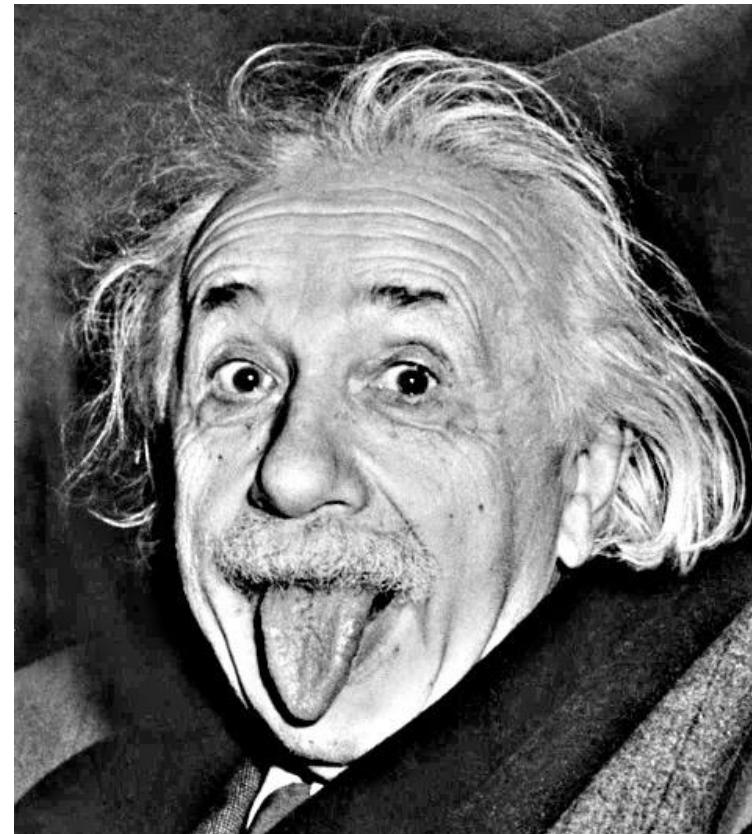
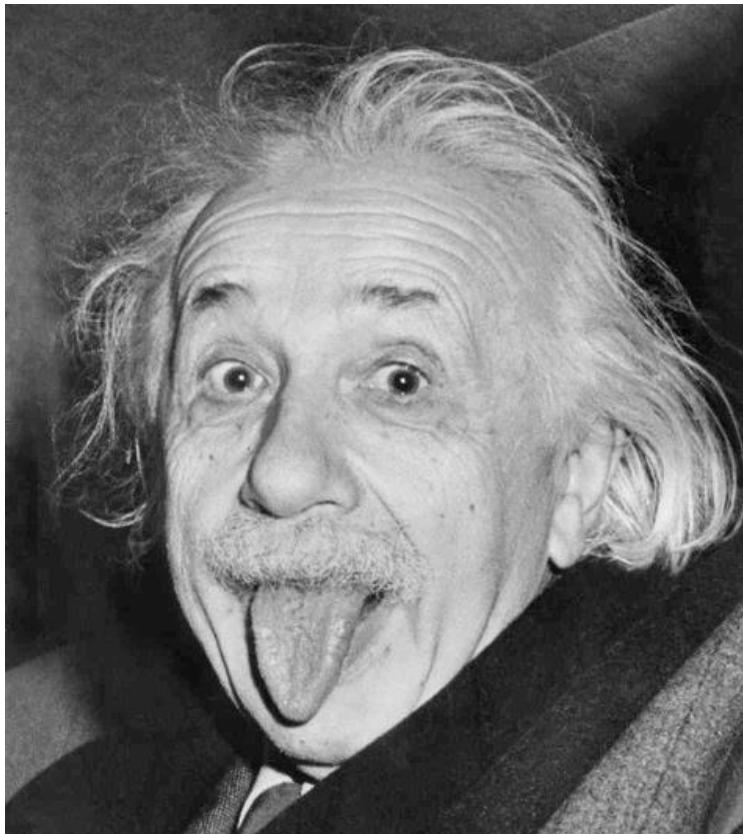


before



after

Sharpening examples



Sharpening examples



Sharpening examples



Sharpening examples



do you see
any problems
in this image?

Do not overdo it with sharpening



original



sharpened



oversharpened

What is wrong in this image?

Convolution

Convolution for 1D continuous signals

Definition of filtering as convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

filtered signal filter input signal notice the flip

Convolution for 1D continuous signals

Definition of filtering as convolution:

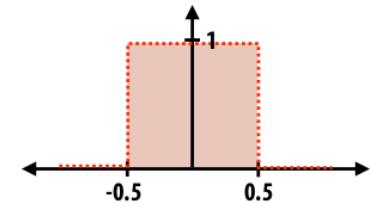
$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

filtered signal filter input signal notice the flip

Consider the box filter example:

1D continuous
box filter

$$f(x) = \begin{cases} 1 & |x| \leq 0.5 \\ 0 & \text{otherwise} \end{cases}$$



filtering output is a
blurred version of
g

$$(f * g)(x) = \int_{-0.5}^{0.5} g(x - y)dy$$

Convolution for 2D discrete signals

Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

filtered image filter input image notice the flip

Convolution for 2D discrete signals

Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

filtered image filter input image notice the flip

If the filter $f(i, j)$ is non-zero only within $-1 \leq i, j \leq 1$, then

$$(f * g)(x, y) = \sum_{i,j=-1}^{1} f(i, j)I(x - i, y - j)$$

The kernel we saw earlier is the 3x3 matrix representation of $f(i, j)$.

Convolution vs correlation

Definition of discrete 2D convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

← notice the flip

Definition of discrete 2D correlation:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x + i, y + j)$$

← notice the lack of a flip

- Most of the time won't matter, because our kernels will be symmetric.
- Will be important when we discuss frequency-domain filtering (lectures 5-6).

Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

$$\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 1 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline \end{array}$$

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

$$\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 1 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline \end{array}$$

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has $M \times M$ pixels and the filter kernel has size $N \times N$:

- What is the cost of convolution with a non-separable filter?

Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

$$\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 1 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline \end{array}$$

row
column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has $M \times M$ pixels and the filter kernel has size $N \times N$:

- What is the cost of convolution with a non-separable filter? \longrightarrow
- What is the cost of convolution with a separable filter?

$$M^2 \times N^2$$

Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

$$\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 1 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline \end{array}$$

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

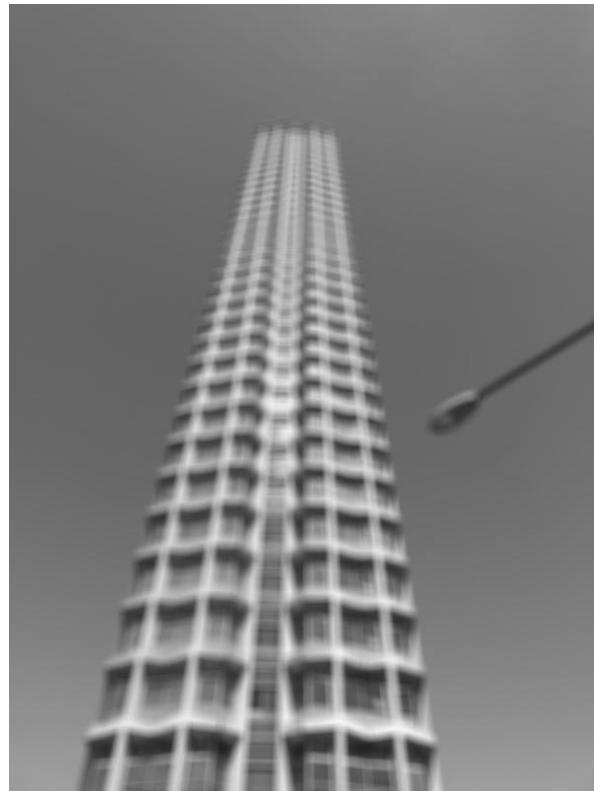
If the image has $M \times M$ pixels and the filter kernel has size $N \times N$:

- What is the cost of convolution with a non-separable filter? $\longrightarrow M^2 \times N^2$
- What is the cost of convolution with a separable filter? $\longrightarrow 2 \times N \times M^2$

A few more filters



original



3x3 box filter

do you see
any problems
in this image?

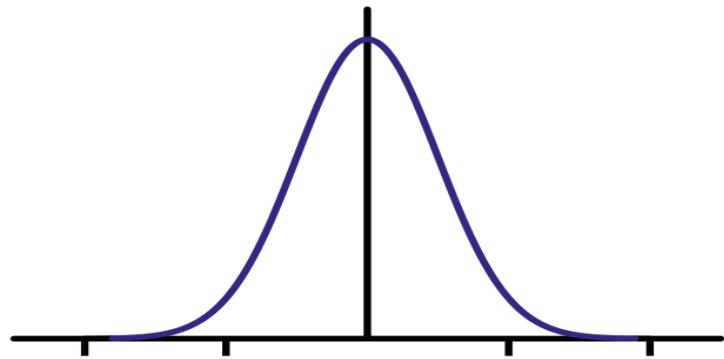
The Gaussian filter

- named (like many other things) after Carl Friedrich Gauss
- kernel values sampled from the 2D Gaussian function:

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

- weight falls off with distance from center pixel
- theoretically infinite, in practice truncated to some maximum distance

Any heuristics for selecting where to truncate?



The Gaussian filter

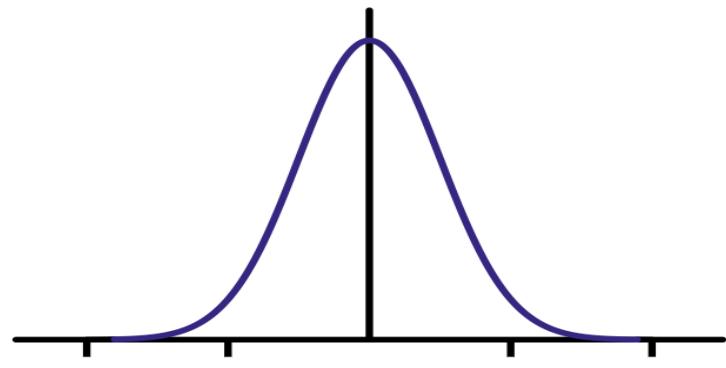
- named (like many other things) after Carl Friedrich Gauss
- kernel values sampled from the 2D Gaussian function:

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

- weight falls off with distance from center pixel
- theoretically infinite, in practice truncated to some maximum distance

Any heuristics for selecting where to truncate?

- usually at $2-3\sigma$



Is this a separable filter?

kernel $\frac{1}{16}$

1	2	1
2	4	2
1	2	1

The Gaussian filter

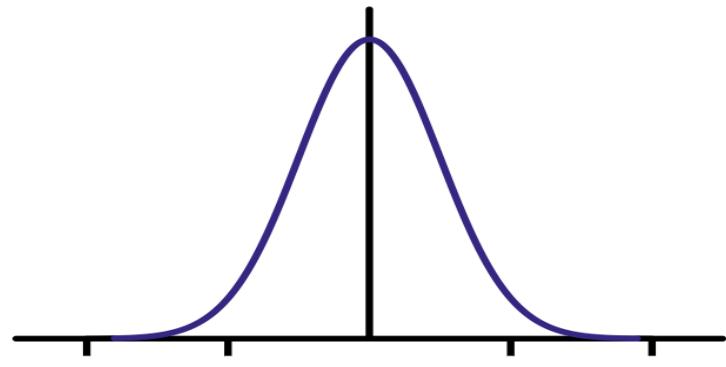
- named (like many other things) after Carl Friedrich Gauss
- kernel values sampled from the 2D Gaussian function:

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

- weight falls off with distance from center pixel
- theoretically infinite, in practice truncated to some maximum distance

Any heuristics for selecting where to truncate?

- usually at $2-3\sigma$

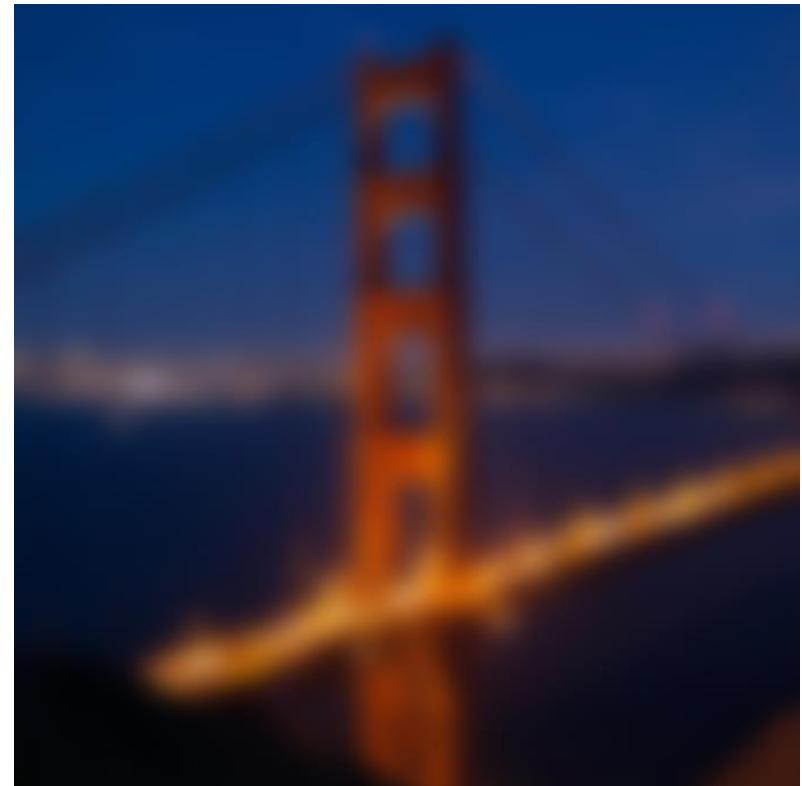


Is this a separable filter? Yes!

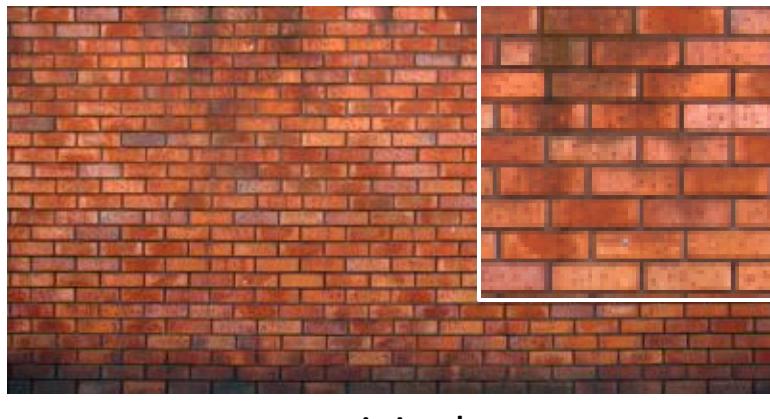
kernel $\frac{1}{16}$

1	2	1
2	4	2
1	2	1

Gaussian filtering example

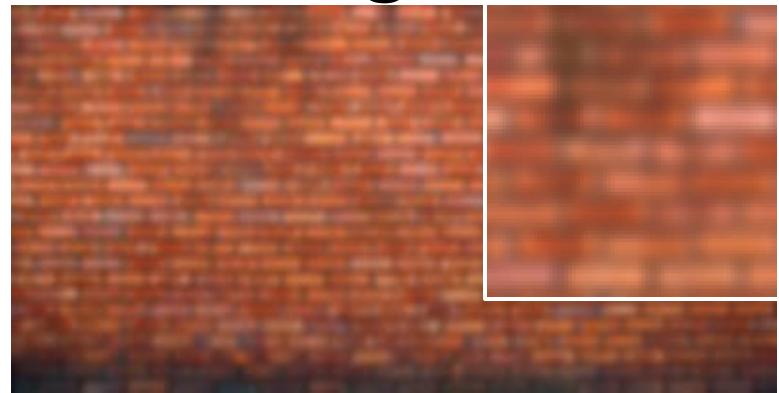


Gaussian vs box filtering

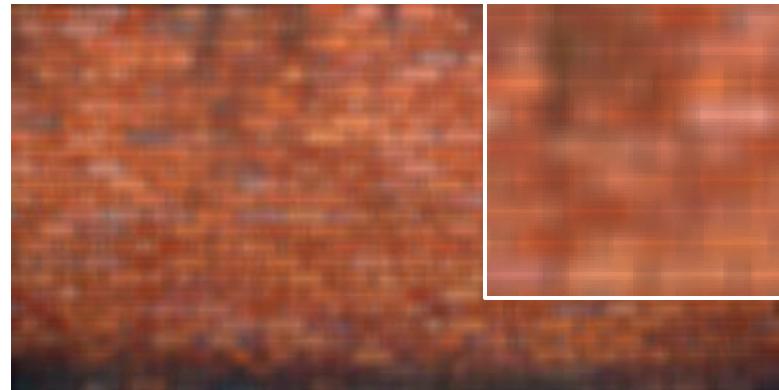


original

Which blur do you like better?



7x7 Gaussian



7x7 box

Correlation and Convolution

- 2d correlation

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

Correlation and Convolution

- 2d correlation

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

- 2d convolution

$$h[m,n] = \sum_{k,l} f[k,l] I[m-k, n-l]$$

e.g., `h = scipy.signal.convolve2d(f, I)`

Convolution is the same as correlation with a 180° rotated filter kernel.
Correlation and convolution are identical when the filter kernel is symmetric.

* Symmetric in the geometric sense, not in the matrix linear algebra sense.

Key properties of linear filters

Linearity:

$$\text{imfilter}(I, f_1 + f_2) = \text{imfilter}(I, f_1) + \text{imfilter}(I, f_2)$$

Shift invariance:

Same behavior given intensities regardless of pixel location m, n

$$\text{imfilter}(I, \text{shift}(f)) = \text{shift}(\text{imfilter}(I, f))$$

Any linear, shift-invariant operator can be represented as a convolution.

Convolution properties

Commutative: $a * b = b * a$

- Conceptually no difference between filter and signal
- But filtering implementations might break this equality, e.g., image edges

Associative: $a * (b * c) = (a * b) * c$

- Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Why important?
- Correlation is not associative (rotation effect)

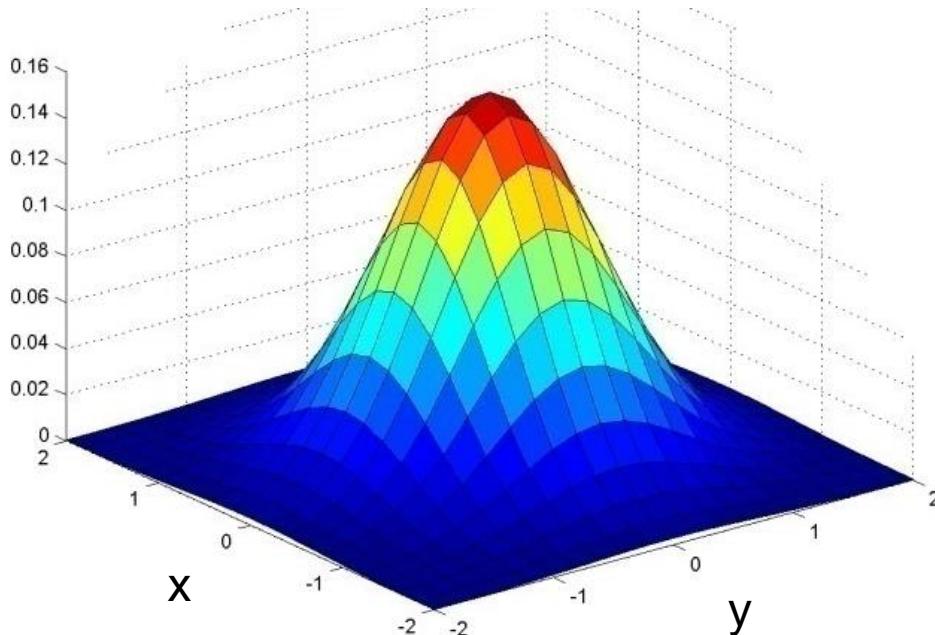
Distributes over addition: $a * (b + c) = (a * b) + (a * c)$

Scalars factor out: $ka * b = a * kb = k(a * b)$

Identity: unit impulse $e = [0, 0, 1, 0, 0]$, $a * e = a$

Important filter: Gaussian

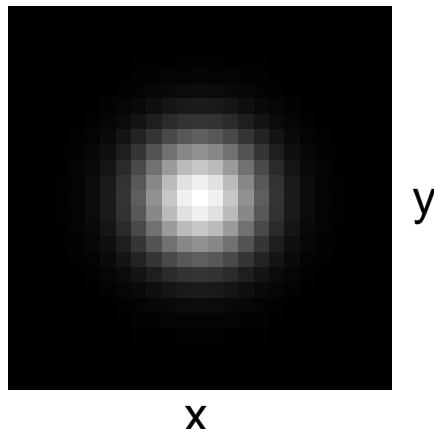
Weight contributions of neighboring pixels by nearness



x	0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013	
0.022	0.097	0.159	0.097	0.022	
0.013	0.059	0.097	0.059	0.013	
0.003	0.013	0.022	0.013	0.003	

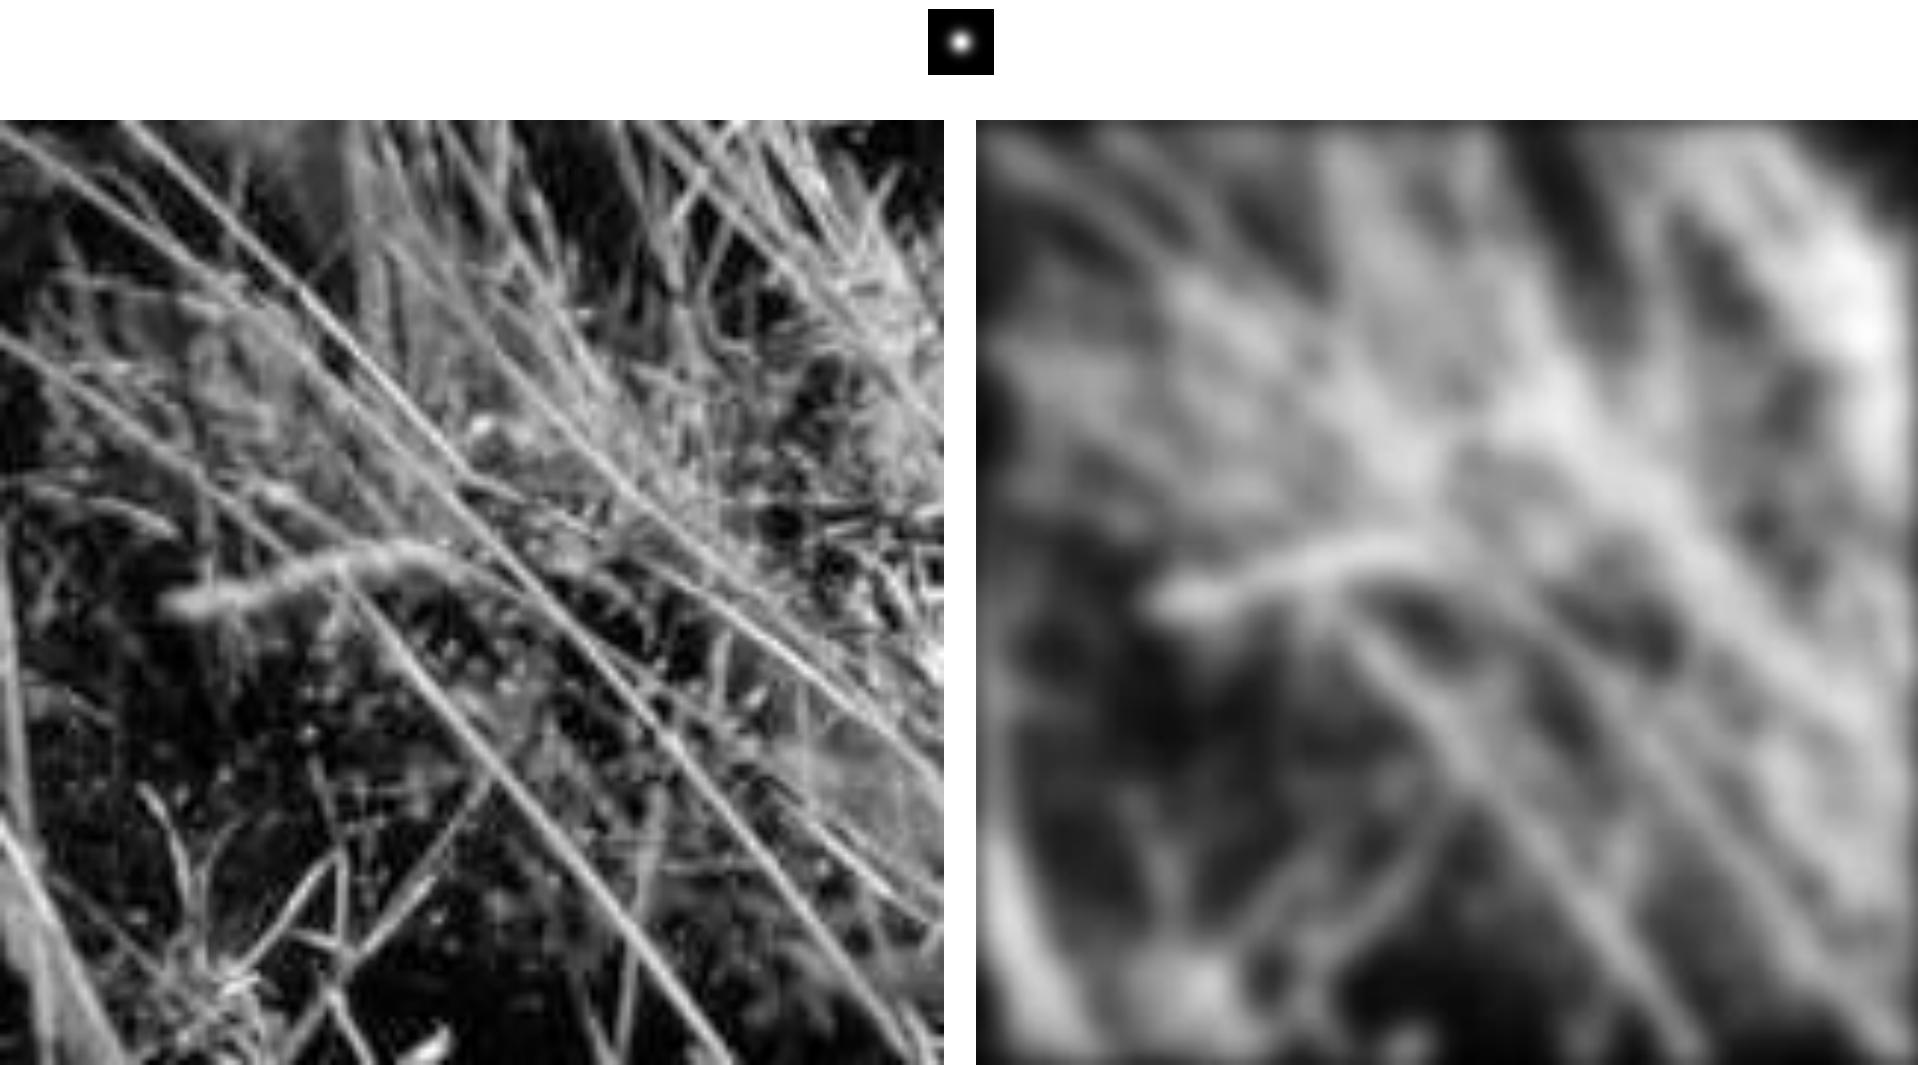
Kernel size 5×5 ,
Standard deviation $\sigma = 1$

Viewed
from top

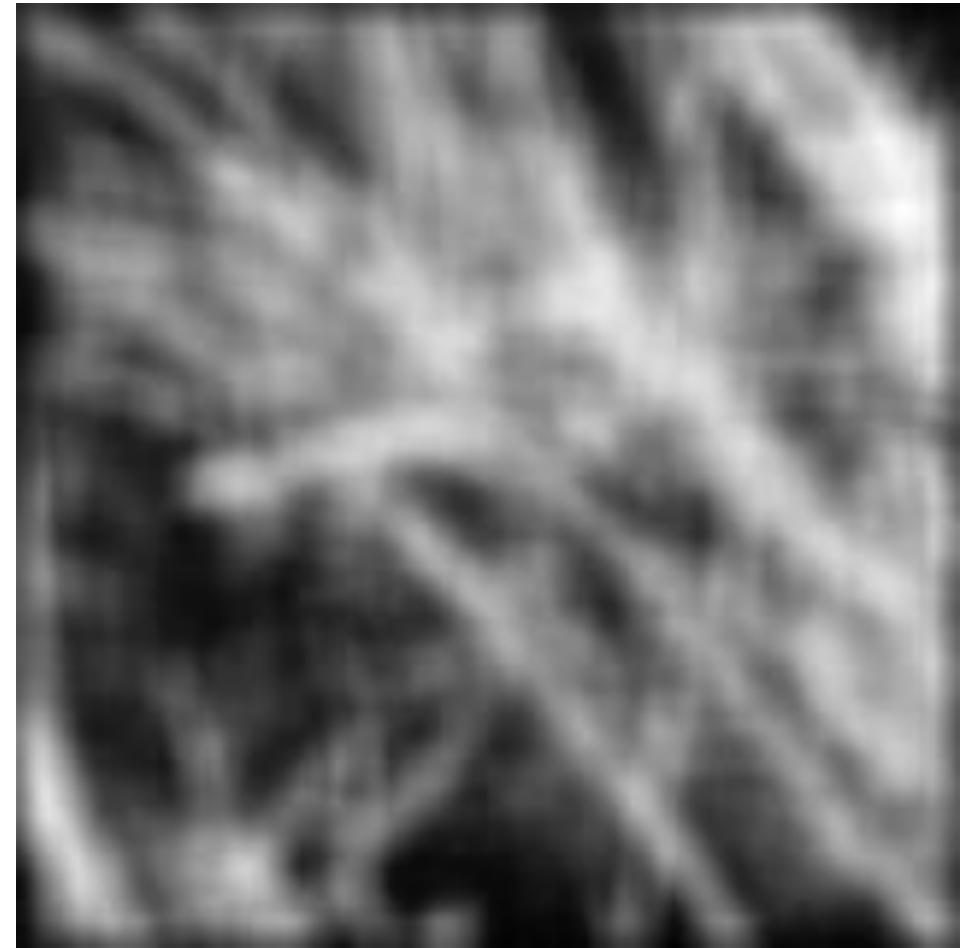


$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Smoothing with Gaussian filter



Smoothing with box filter



Gaussian filters

Gaussian convolved with Gaussian...

...is another Gaussian

- So can smooth with small-width kernel, repeat, and get same result as larger-width kernel
- Convolving twice with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$

Separable kernel

- Factors into product of two 1D Gaussians

Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

2D convolution
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

Perform convolution
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} 11 & & \\ 18 & & \\ 18 & & \end{matrix}$$

Followed by convolution
along the remaining column:

Separability

Why is separability useful in practice?

MxN image, PxQ filter

- 2D convolution: $\sim MNPQ$ multiply-adds
- Separable 2D: $\sim MN(P+Q)$ multiply-adds

Speed up = $PQ/(P+Q)$

9x9 filter = $\sim 4.5x$ faster

Practical matters

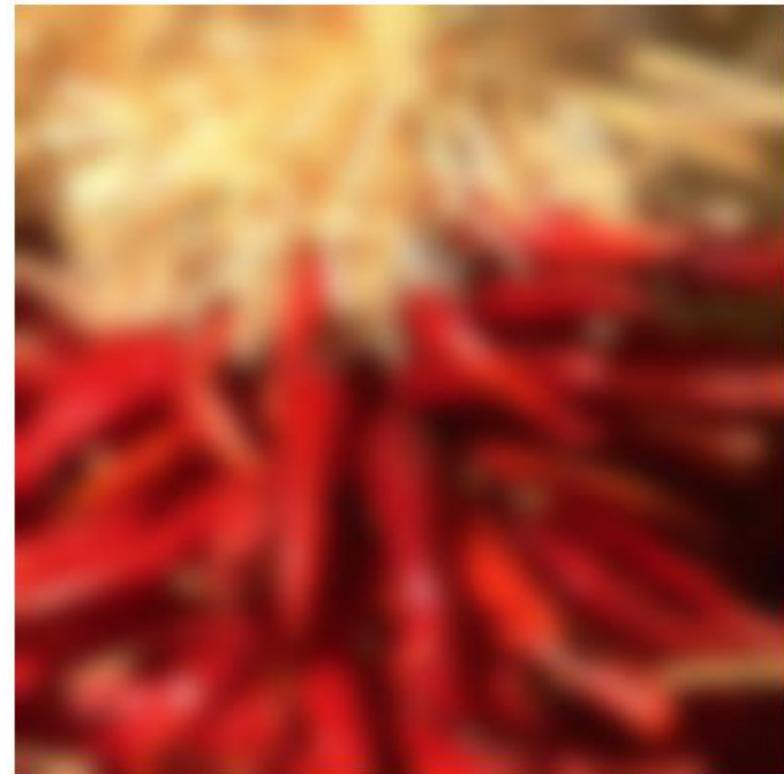
How big should the filter be?

- Values at edges should be near zero
- Gaussians have infinite extent...
- Rule of thumb for Gaussian: set filter half-width to about $3\ \sigma$

Practical matters

What about near the edge?

- The filter window falls off the edge of the image
- Need to extrapolate
- methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



Convolution in Convolutional Neural Networks

- Convolution is the basic operation in CNNs
- Learning convolution kernels allows us to learn which ‘features’ provide useful information in images.

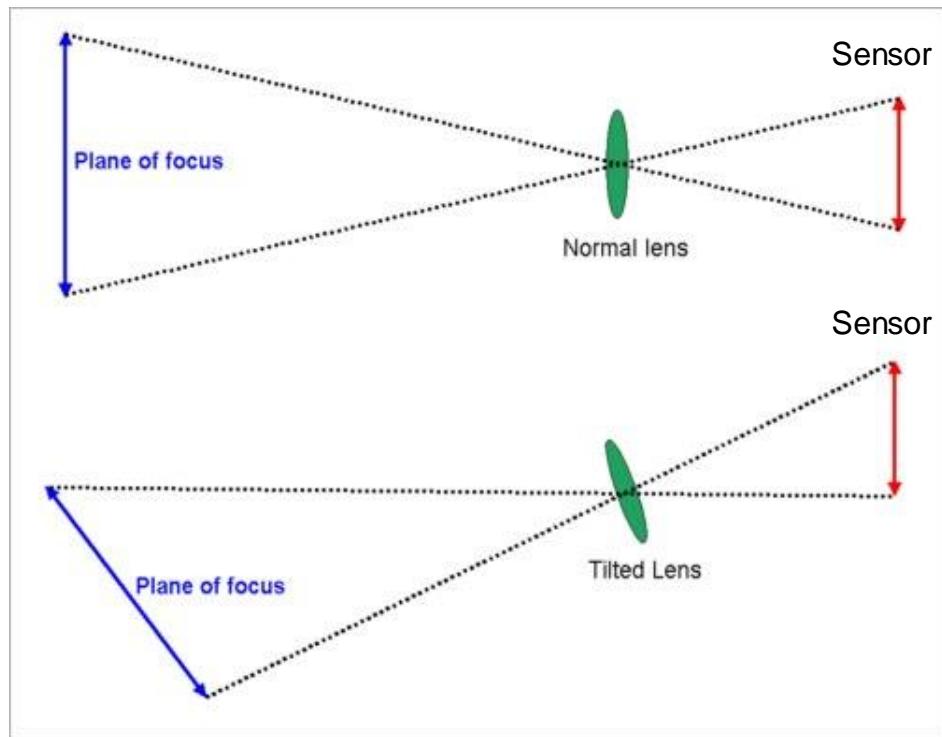
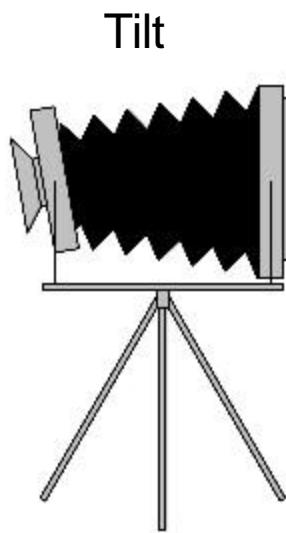
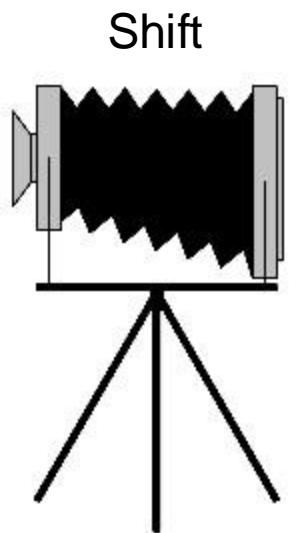


Ben Thomas

Tilt-shift photography



Tilt shift camera



Can we fake tilt shift?

We need to blur the image

- OK, now we know how to do that.

Can we fake tilt shift?

We need to blur the image

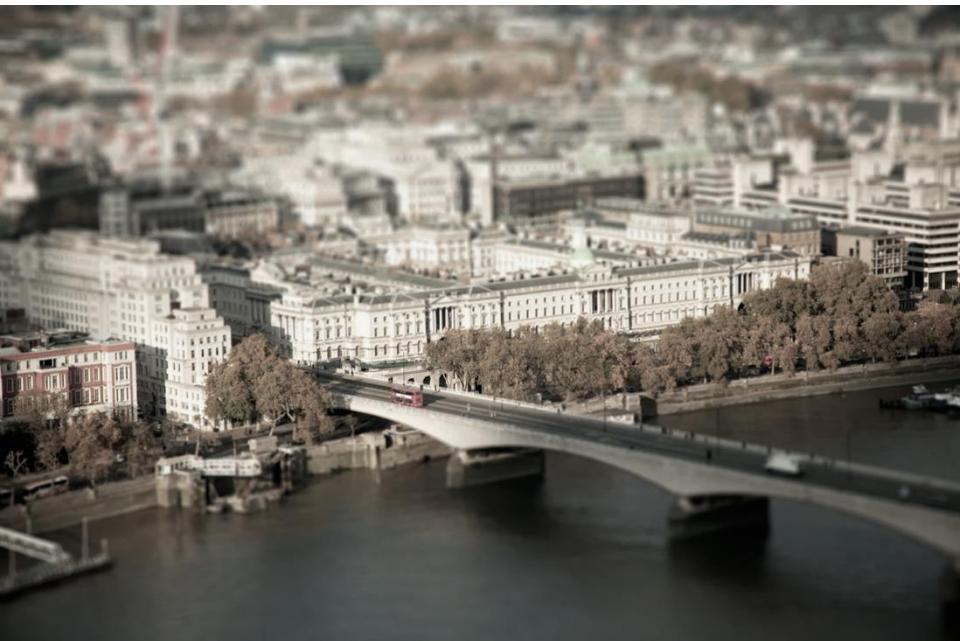
– OK, now we know how to do that.

We need to blur progressively more away from our ‘fake’ focal point



But can I make it look more like a toy?

- Boost *saturation* – toys are very colorful
- We'll learn how to do this when we discuss color

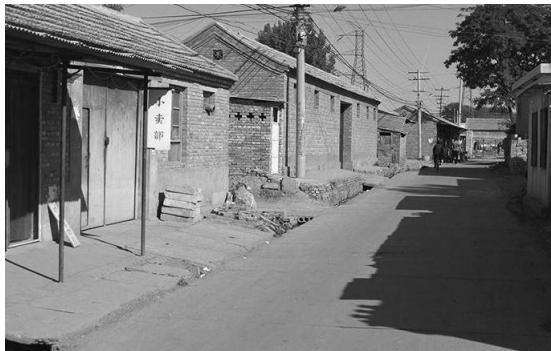


Think-Pair-Share

* = Convolution operator

$$\begin{array}{l} 1) \quad \underline{\quad} = D * B \\ 2) \quad \overline{A} = \underline{\quad} * \underline{\quad} \\ 3) \quad F = \overline{D} * \underline{\quad} \\ 4) \quad \underline{\quad} = D * \overline{D} \end{array}$$

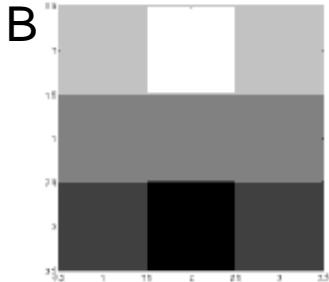
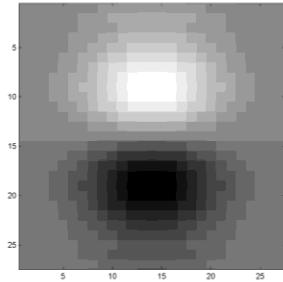
D



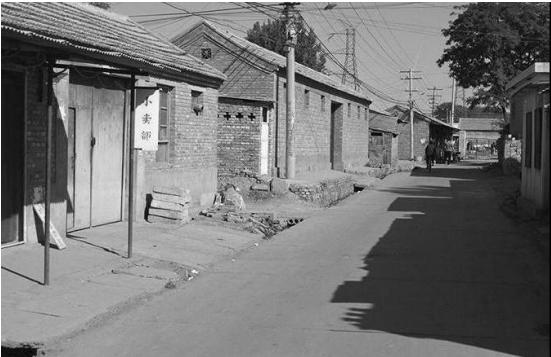
H



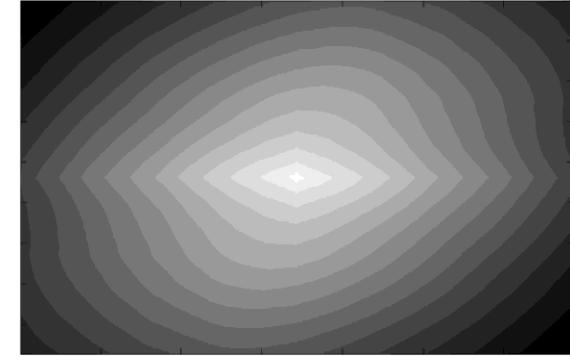
A



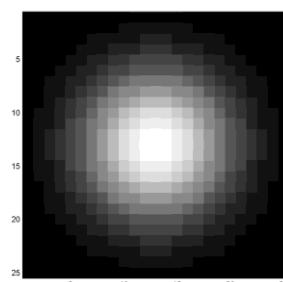
F



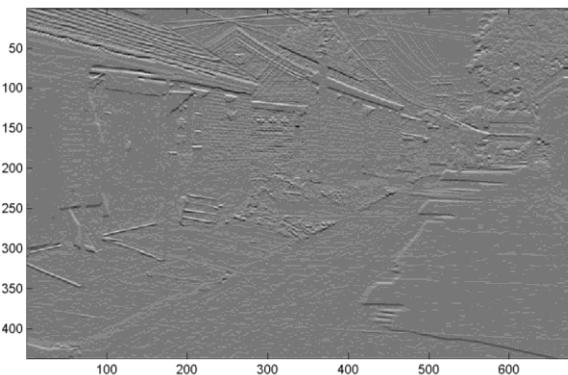
I



C



G

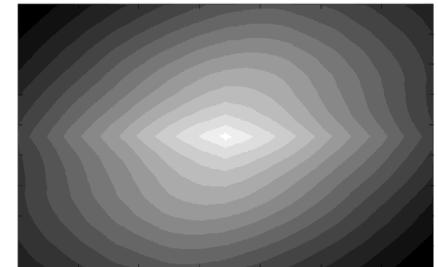


$$I = D * D$$

D (275 x 175 pixels)



I (from slide – 275 x 175)



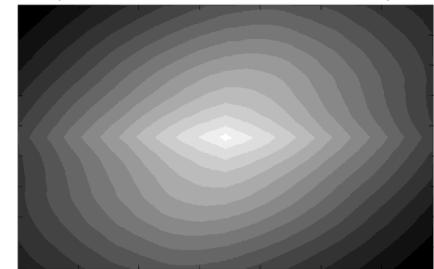
“...something to do with lack of content (black) at edges...”

$$I = D * D$$

D (275 x 175 pixels)



I (from slide – 275 x 175)



```
>> D = img_to_float32( io.imread( 'convexample.png' ) )
```

```
>> I = convolve2d( D, D )
```

I_norm

```
>> np.max(I)
```

```
1.1021e+04
```

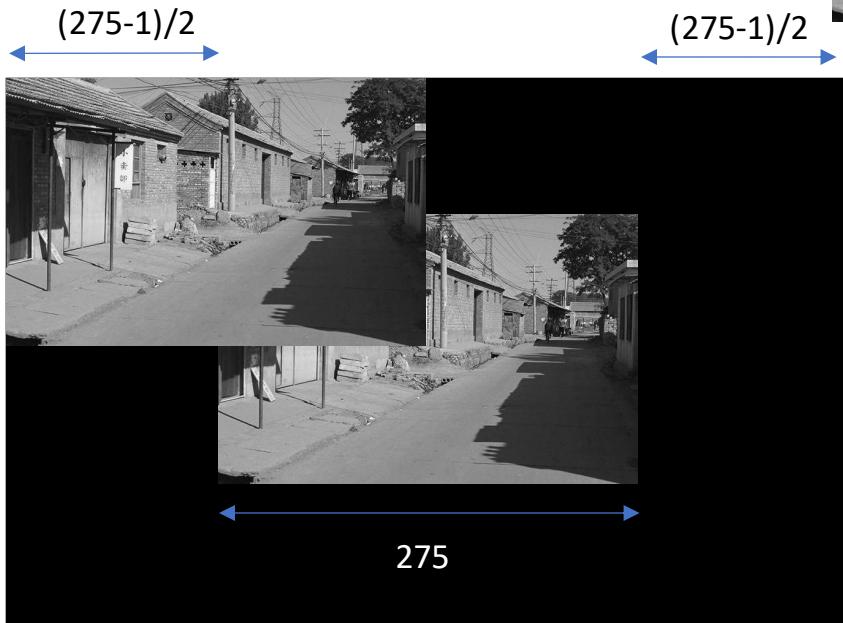
```
# Normalize for visualization
```

```
>> I_norm = (I - np.min(I)) / (np.ma
```

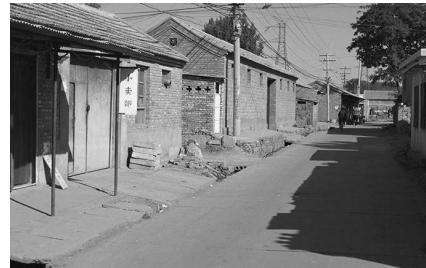
```
>> plt.imshow( I_norm )
```



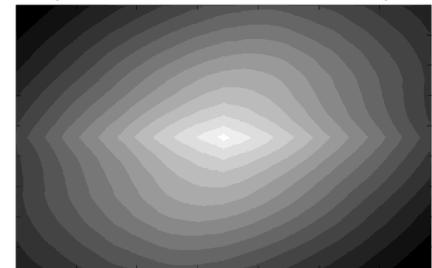
$$I = D * D$$



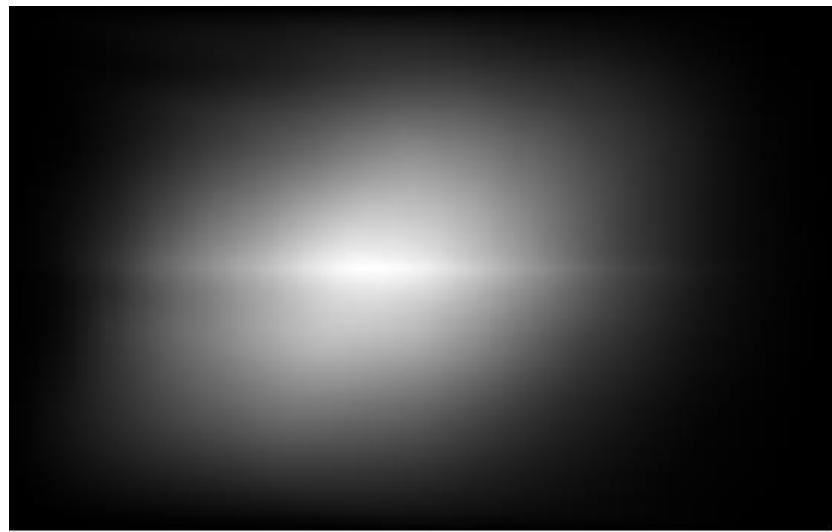
D (275 x 175 pixels)



I (from slide – 275 x 175)



I_norm (549 x 349 pixels)



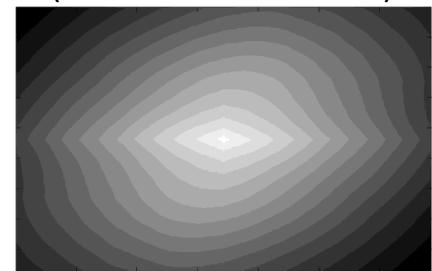
$$\begin{aligned} \text{For } x: 275 + (275-1)/2 + (275-1)/2 \\ = 549 \end{aligned}$$

$$I = D * D$$

D (275 x 175 pixels)

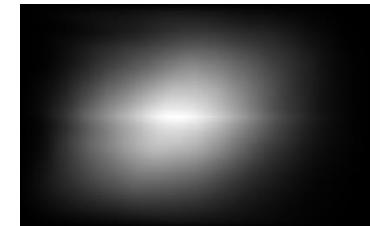


I (from slide – 275 x 175)



>> I = convolve2d(D, D, mode='full')

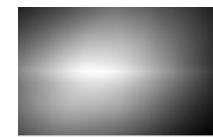
(Default; pad with zeros)



549 x 349

>> I = convolve2d(D, D, mode='same')

(Return same size as D)



275 x 175

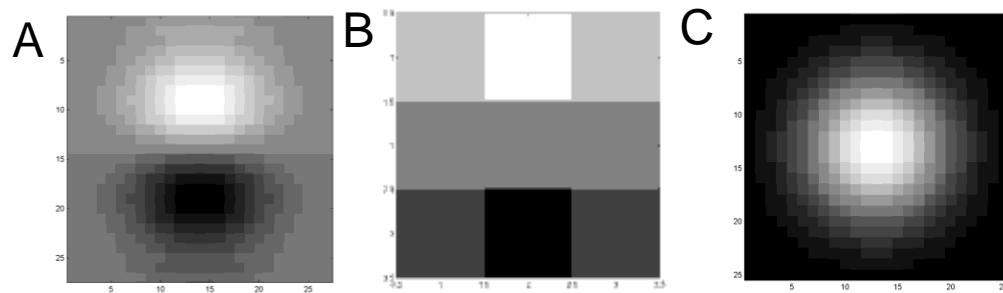
>> I = convolve2d(D, D, mode='valid')

(No padding)



Value = 10528.3
1x1

$A = B * C$ – “because it kind of looks like it.”



C is a Gaussian filter
(or something close to it),
and we know that it ‘blurs’.

When the filter ‘looks like’ the image = ‘template matching’

Filtering viewed as comparing an image of
what you want to find against all image regions.

For symmetric filters: use either convolution or correlation.

For nonsymmetric filters: correlation is template matching.

Filtering: Correlation and Convolution

- 2d correlation

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

- 2d convolution

$$h[m,n] = \sum_{k,l} f[k,l] I[m-k, n-l]$$

e.g., `h = scipy.signal.convolve2d(f, I)`

Convolution is the same as correlation with a 180° rotated filter kernel.
Correlation and convolution are identical when the filter kernel is symmetric.

OK, so let's test this idea. Let's see if we can use correlation to 'find' the parts of the image that look like the filter.

```
>> f = D[ 57:117, 107:167 ]
```

Expect response 'peak' in middle of I

```
>> I = correlate2d( D, f, 'same' )
```

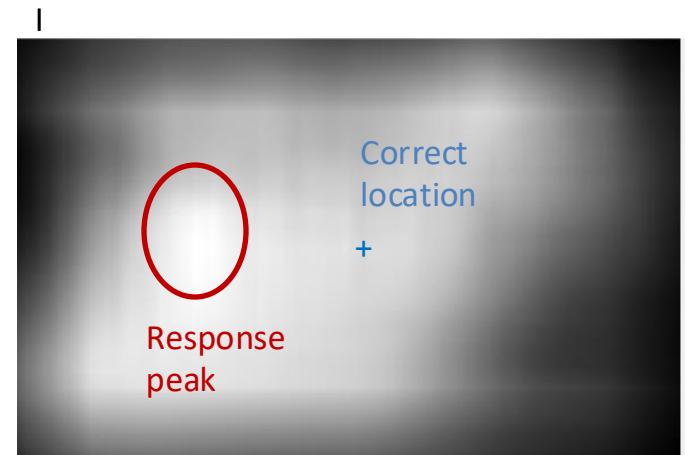
Hmm...

That didn't work – why not?

D (275 x 175 pixels)



f
61 x 61



Correlation

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

As brightness in I increases, the response in h will increase, as long as f is positive.

Overall brighter regions will give higher correlation response -> not useful!

OK, so let's subtract the mean

```
>> f = D[ 57:117, 107:167 ]  
>> f2 = f - np.mean(f)  
>> D2 = D - np.mean(D)
```



D2 (275 x 175 pixels)



*Remember –
float data type –
goes negative!*

*Now zero centered.
Score is higher only when dark parts
match and when light parts match.*

```
>> I2 = correlate2d( D2, f2, 'same' )
```



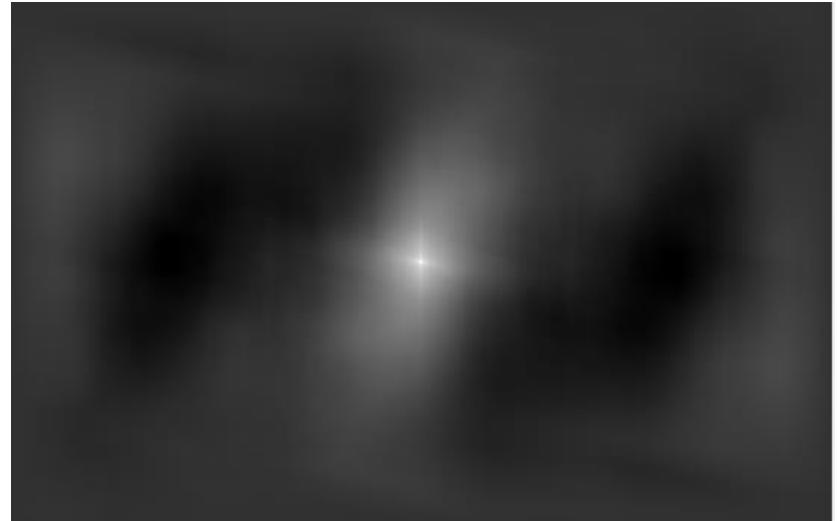
D2 (275 x 175 pixels)



Or our original example: D2 correlated with itself

```
>> I3 = correlate2d( D2, D2, 'full' )
```

I3



What happens with convolution?

```
>> f = D[ 57:117, 107:167 ]
```

```
>> f2 = f - np.mean(f)
```

```
>> D2 = D - np.mean(D)
```

```
>> I2 = convolve2d( D2, f2, 'same' )
```

D2 (275 x 175 pixels)



f2
61 x 61

I2



NON-LINEAR FILTERS

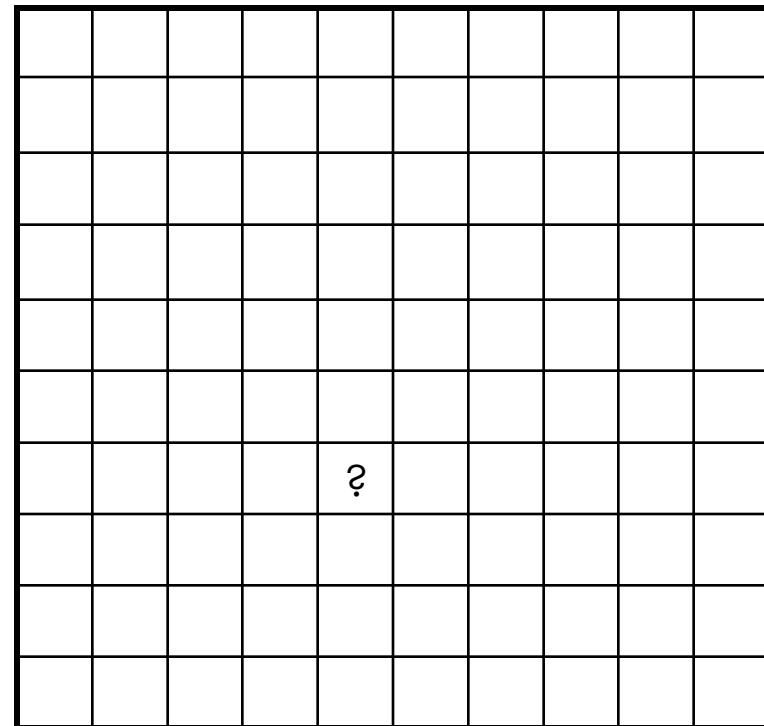
Median filters

- Operates over a window by selecting the median intensity in the window.
- ‘Rank’ filter as based on ordering of gray levels
 - E.G., min, max, range filters

Median filter?

$$I[\cdot, \cdot]$$

$h[.,.]$



Median filters

- Operates over a window by selecting the middle intensity in the window.
- What advantage does a median filter have over a mean filter?

Noise – Salt and Pepper Jack



Mean Jack – 3 x 3 filter



Very Mean Jack – 11 x 11 filter



Noisy – Salt and Pepper Jack



Median Jack – 3 x 3



Very Median Jack – 11 x 11



Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?

Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?

Interpretation: Median filtering is sorting.

Review: questions

1. Write down a 3x3 filter that both:
 - Returns a positive value if the average value of the 4-adjacent neighbors is less than the center,
 - Returns a negative value otherwise.
2. Write down a filter that will compute the gradient in the x-direction:

$$\text{grad}_x(x, y) = \text{im}(x+1, y) - \text{im}(x, y) \text{ for each } x, y$$

Review: questions

1. Write down a 3x3 filter that both:

- Returns a positive value if the average value of the 4-adjacent neighbors is less than the center,
 - Returns a negative value otherwise.
- $$\begin{bmatrix} 0 & -1/4 & 0; \\ -1/4 & 1 & -1/4; \\ 0 & -1/4 & 0 \end{bmatrix}$$

2. Write down a filter that will compute the gradient in the x-direction:

$$\text{grad}_x(x, y) = \text{im}(x+1, y) - \text{im}(x, y) \text{ for each } x, y$$

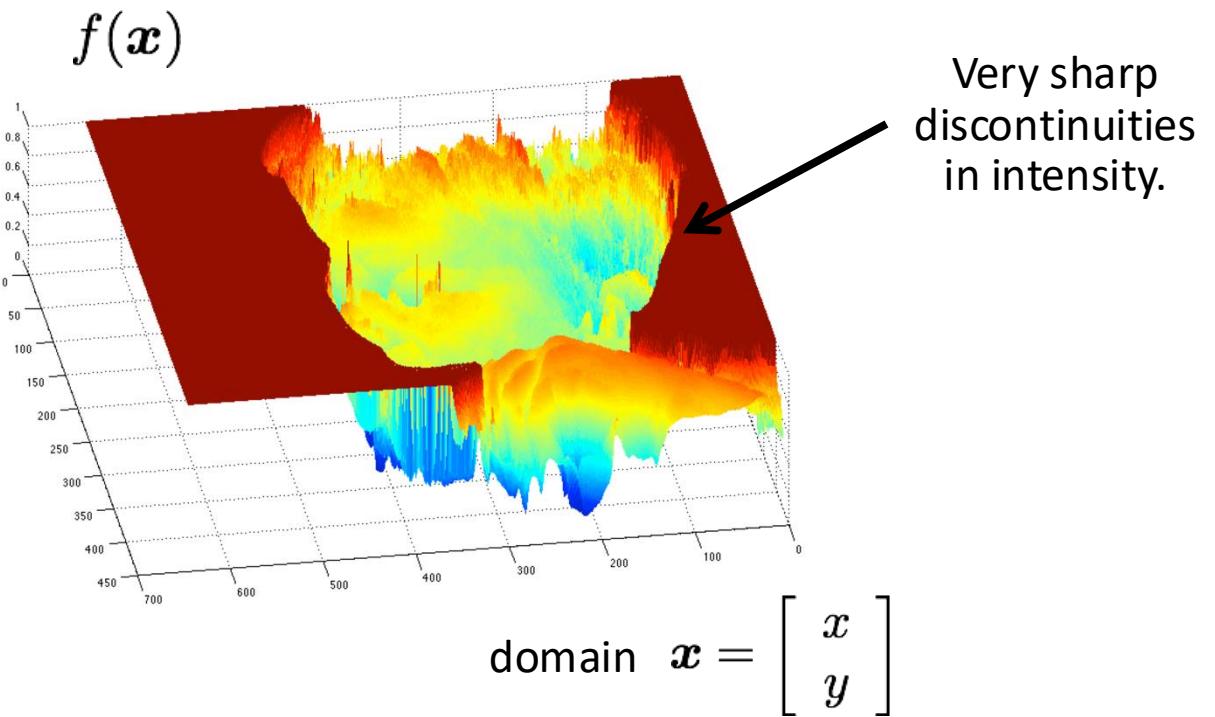
$$\begin{bmatrix} 0 & -1 & 1 \end{bmatrix}$$

Image gradients

What are image edges?



grayscale image



Detecting edges

How would you go about detecting edges in an image (i.e., discontinuities in a function)?

Detecting edges

How would you go about detecting edges in an image (i.e., discontinuities in a function)?

- ✓ You take derivatives: derivatives are large at discontinuities.

How do you differentiate a discrete image (or any other discrete signal)?

Detecting edges

How would you go about detecting edges in an image (i.e., discontinuities in a function)?

- ✓ You take derivatives: derivatives are large at discontinuities.

How do you differentiate a discrete image (or any other discrete signal)?

- ✓ You use finite differences.

Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$

For discrete signals: Remove limit and set $h = 2$

$$f'(x) = \frac{f(x + 1) - f(x - 1)}{2}$$

What convolution kernel does this correspond to?

Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$

For discrete signals: Remove limit and set $h = 2$

$$f'(x) = \frac{f(x + 1) - f(x - 1)}{2}$$

1D derivative filter

1	0	-1
---	---	----

The Sobel filter

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

Sobel filter

What filter
is this?

1D derivative
filter

The Sobel filter

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \text{Sobel filter} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \text{Blurring} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix} \text{1D derivative filter}$$

In a 2D image, does this filter responses along horizontal or vertical lines?

The Sobel filter

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \text{Sobel filter} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \text{Blurring} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix} \text{1D derivative filter}$$

Does this filter return large responses on vertical or horizontal lines?

The Sobel filter

Horizontal Sober filter:

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

What does the vertical Sobel filter look like?

The Sobel filter

Horizontal Sober filter:

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

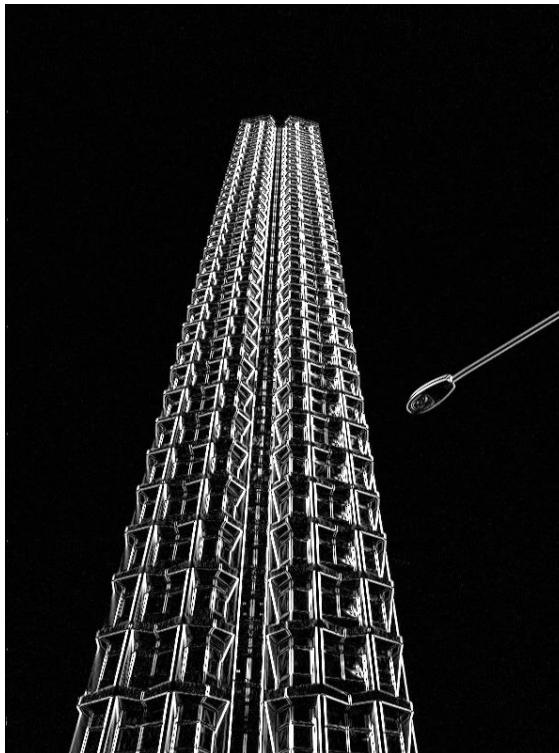
Vertical Sobel filter:

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Sobel filter example



original



which Sobel filter?

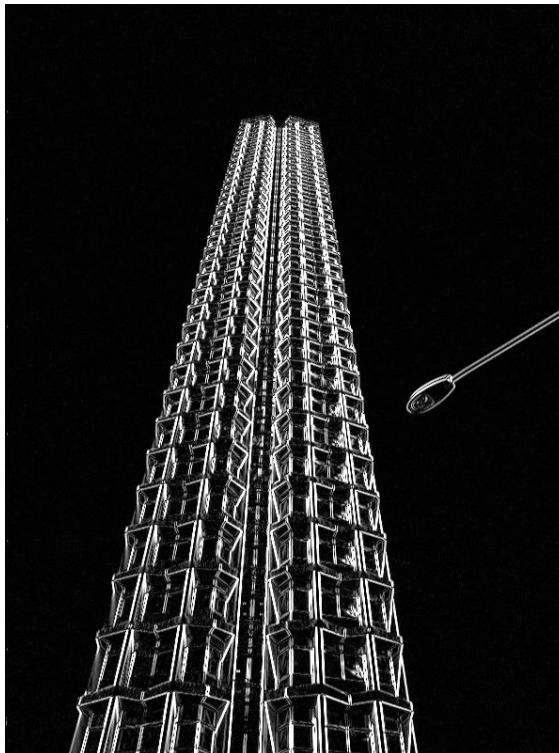


which Sobel filter?

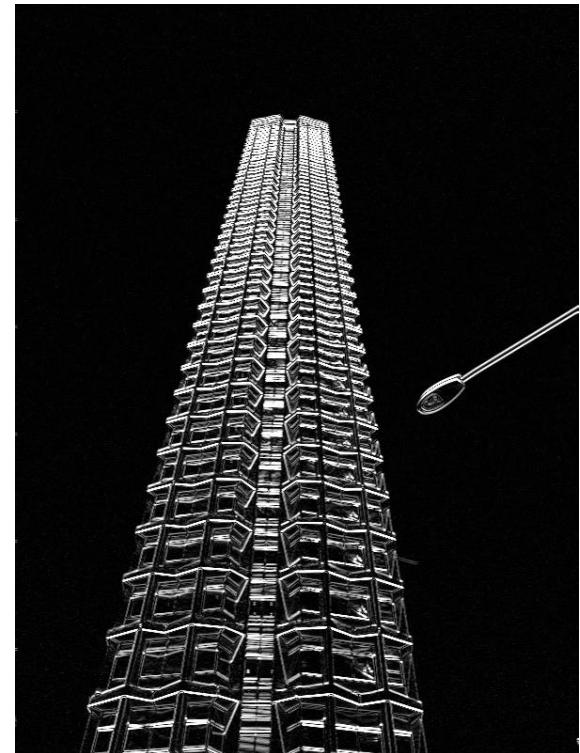
Sobel filter example



original

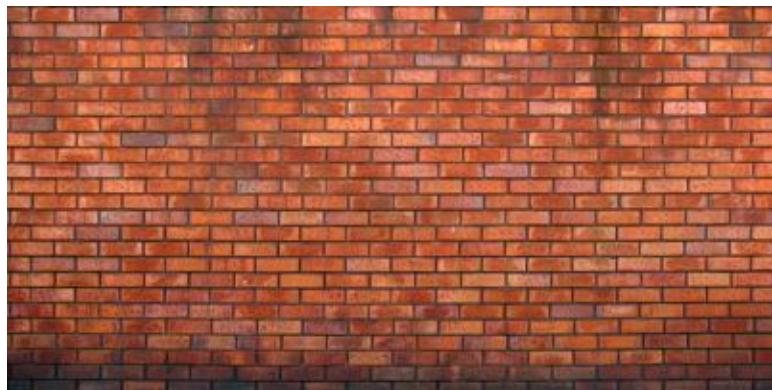


horizontal Sobel filter



vertical Sobel filter

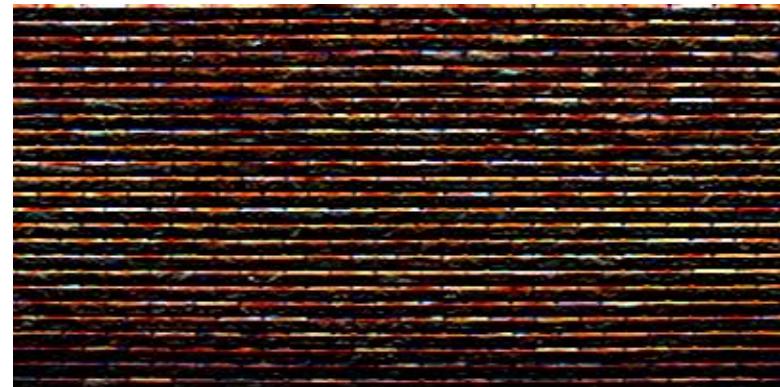
Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter

Several derivative filters

Sobel

1	0	-1	1	2	1
2	0	-2	0	0	0
1	0	-1	-1	-2	-1

Scharr

3	0	-3	3	10	3
10	0	-10	0	0	0
3	0	-3	-3	-10	-3

Prewitt

1	0	-1	1	1	1
1	0	-1	0	0	0
1	0	-1	-1	-1	-1

Roberts

0	1
-1	0

1	0
0	-1

- How are the other filters derived and how do they relate to the Sobel filter?
- How would you derive a derivative filter that is larger than 3x3?

Computing image gradients

1. Select your favorite derivative filters.

$$S_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$S_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

Computing image gradients

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

3. Form the image gradient, and compute its direction and amplitude.

$$\nabla \mathbf{f} = \left[\frac{\partial \mathbf{f}}{\partial x}, \frac{\partial \mathbf{f}}{\partial y} \right]$$

gradient

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

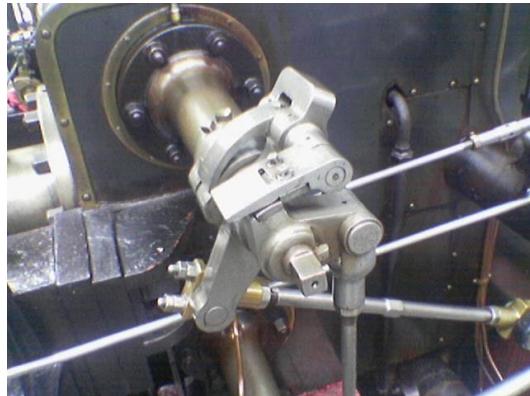
direction

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

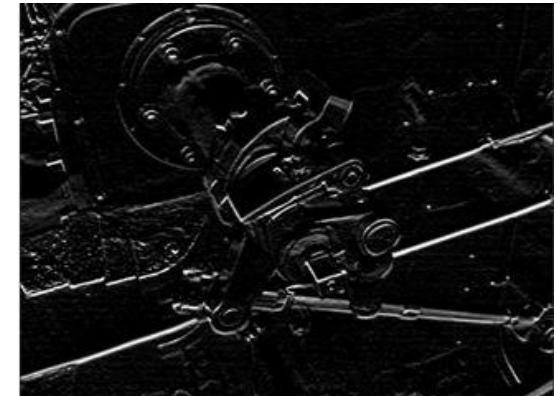
amplitude

Image gradient example

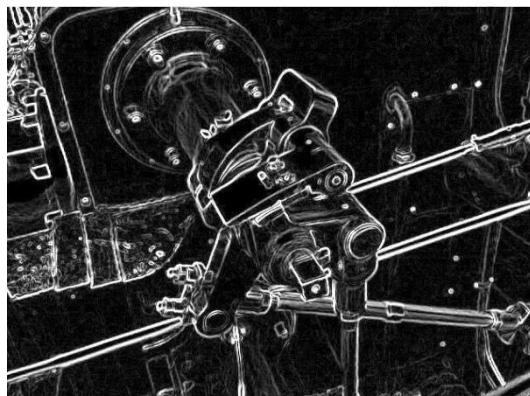
original



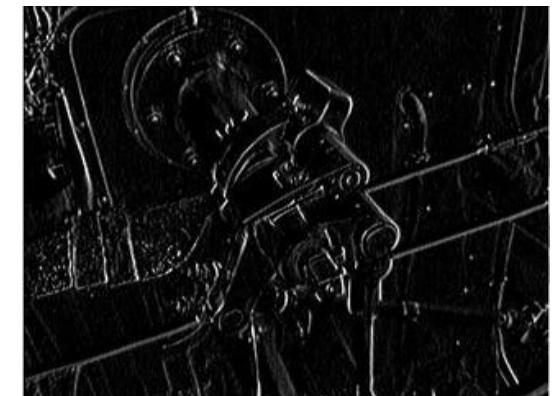
vertical derivative



gradient amplitude



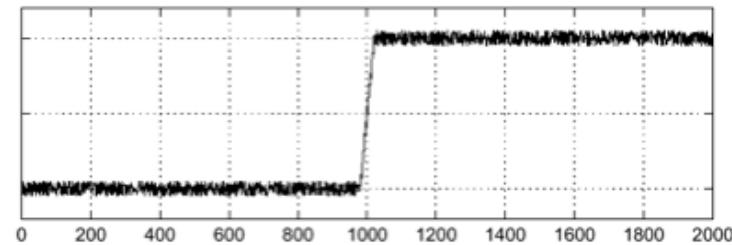
horizontal derivative



How does the gradient direction relate to these edges?

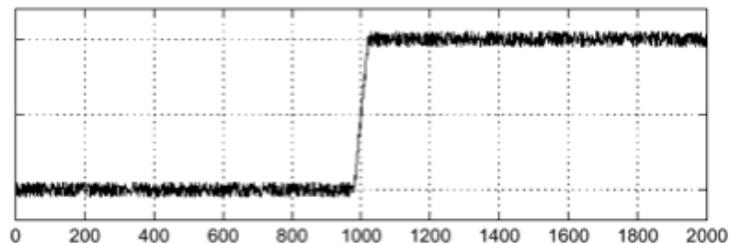
How do you find the edge of this signal?

intensity plot



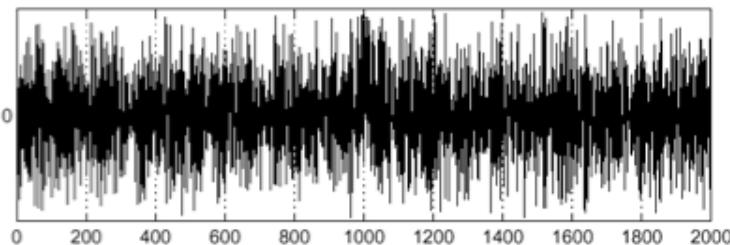
How do you find the edge of this signal?

intensity plot



Using a derivative filter:

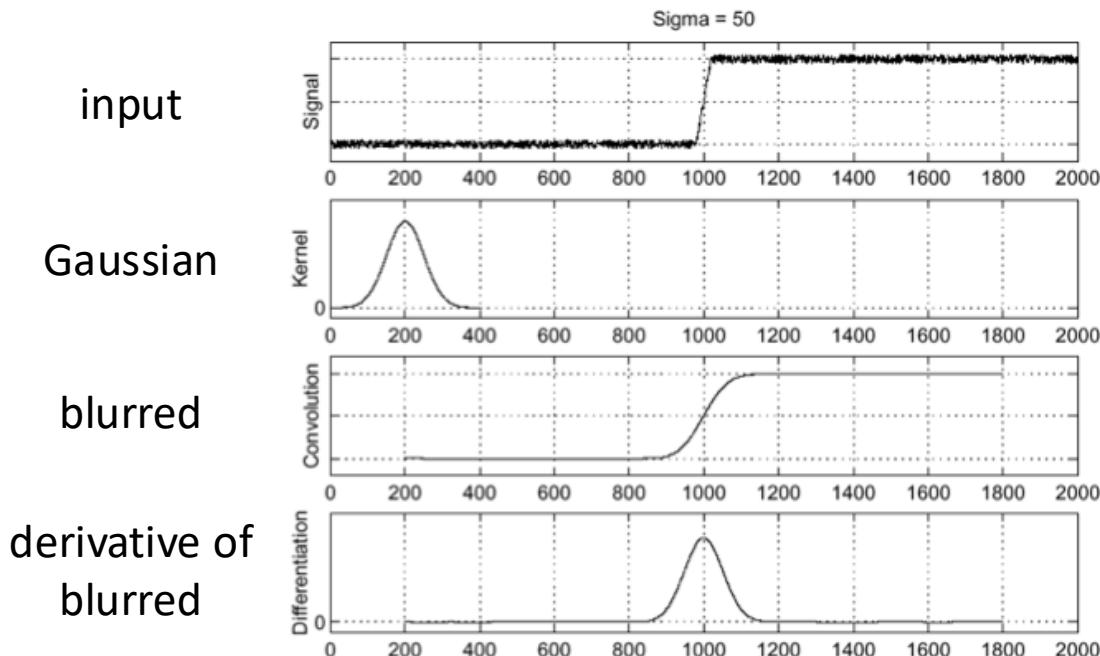
derivative plot



What's the
problem here?

Differentiation is very sensitive to noise

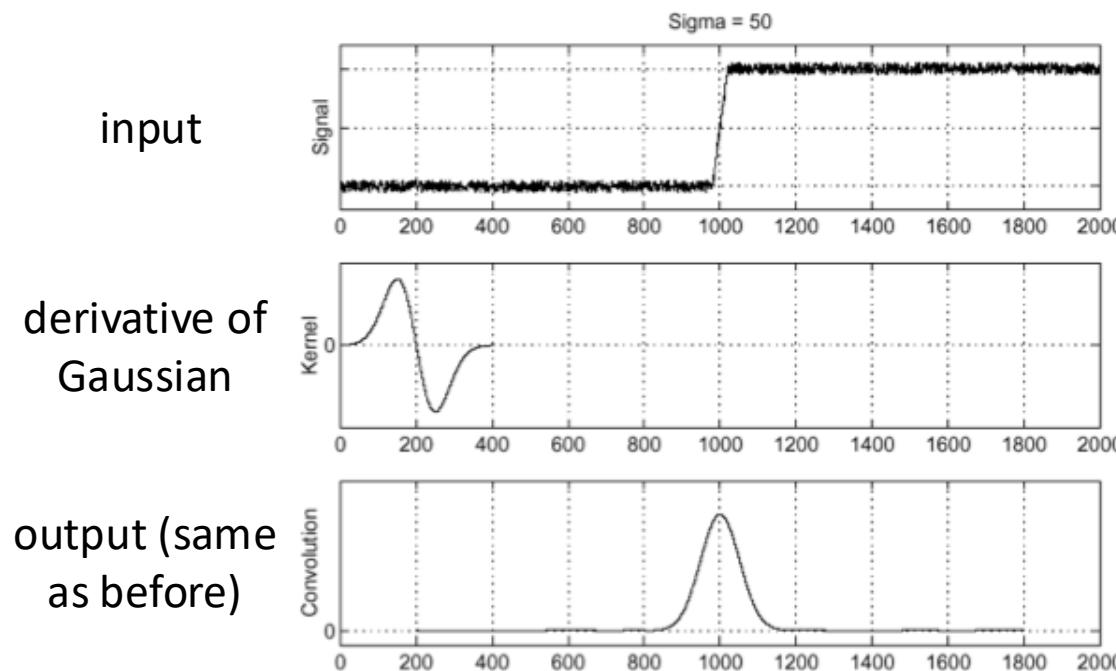
When using derivative filters, it is critical to blur first!



How much
should we blur?

Derivative of Gaussian (DoG) filter

Derivative theorem of convolution: $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$



- How many operations did we save?
- Any other advantages beyond efficiency?

Laplace filter

Basically a second derivative filter.

- We can use finite differences to derive it, as with first derivative filter.

first-order
finite difference $f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$ \rightarrow 1D derivative filter

1	0	-1
---	---	----

second-order
finite difference $f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$ \rightarrow Laplace filter
?

Laplace filter

Basically a second derivative filter.

- We can use finite differences to derive it, as with first derivative filter.

first-order
finite difference $f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$ 1D derivative filter

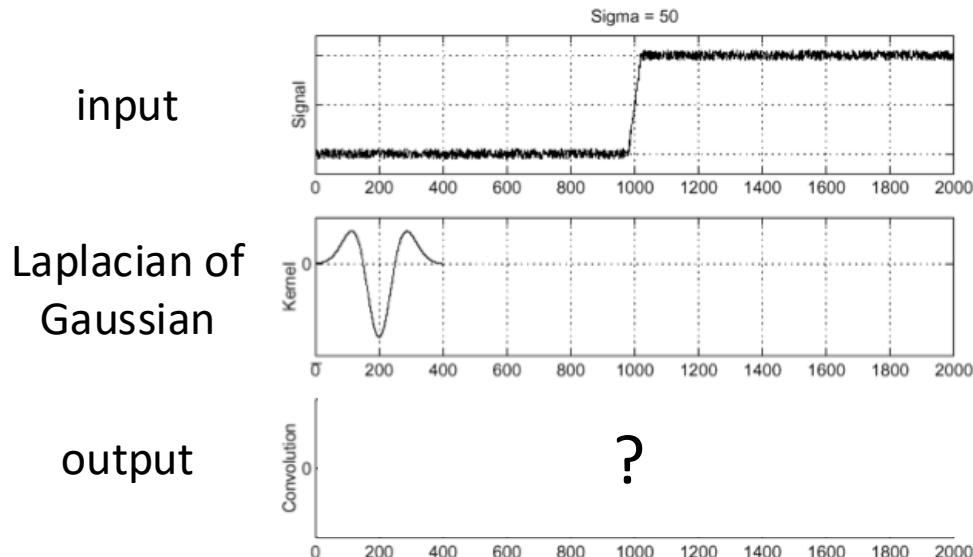
1	0	-1
---	---	----

second-order
finite difference $f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$ Laplace filter

1	-2	1
---	----	---

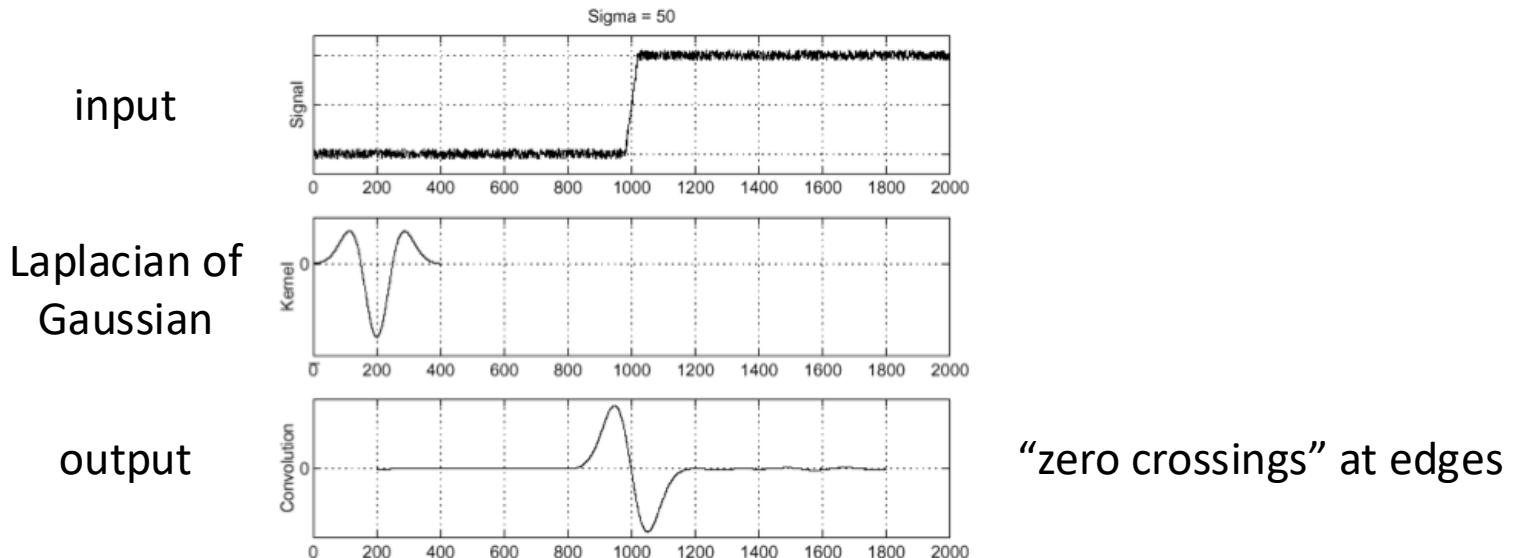
Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering



Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering



Laplace and LoG filtering examples



Laplacian of Gaussian filtering



Laplace filtering

Laplacian of Gaussian vs Derivative of Gaussian

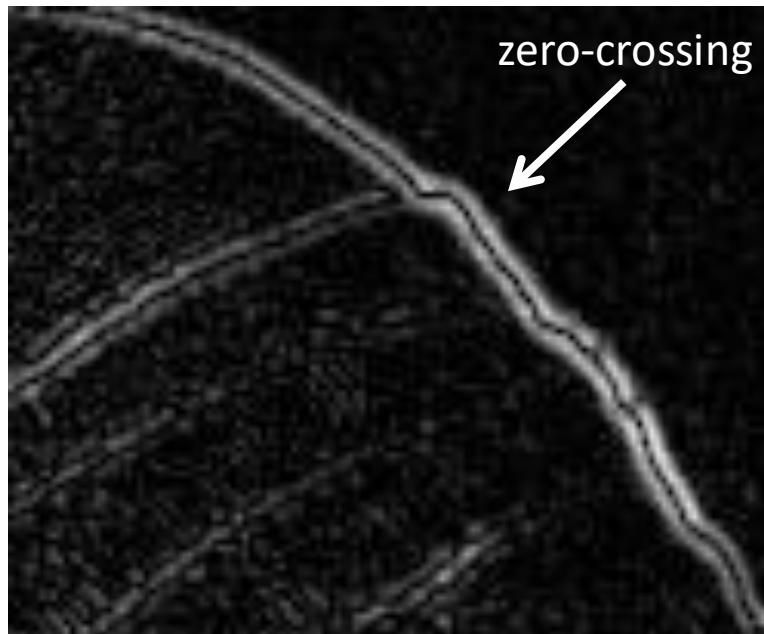


Laplacian of Gaussian filtering

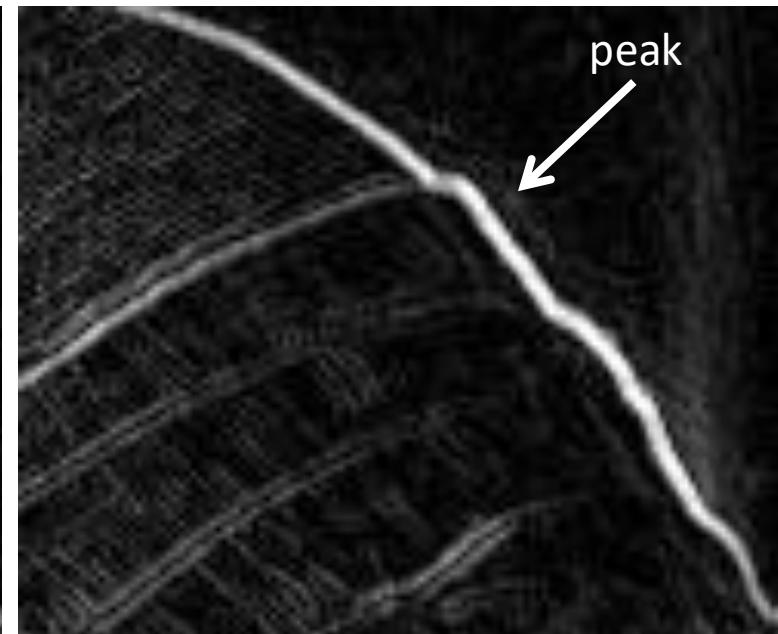


Derivative of Gaussian filtering

Laplacian of Gaussian vs Derivative of Gaussian



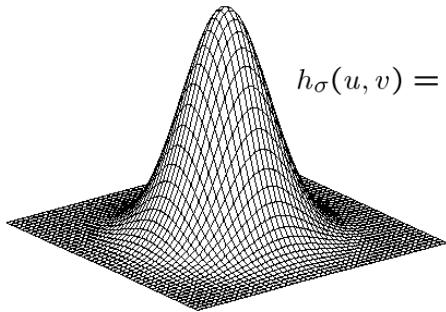
Laplacian of Gaussian filtering



Derivative of Gaussian filtering

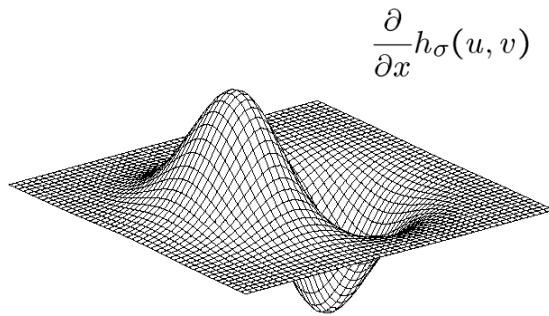
Zero crossings are more accurate at localizing edges
(but not very convenient).

2D Gaussian filters



Gaussian

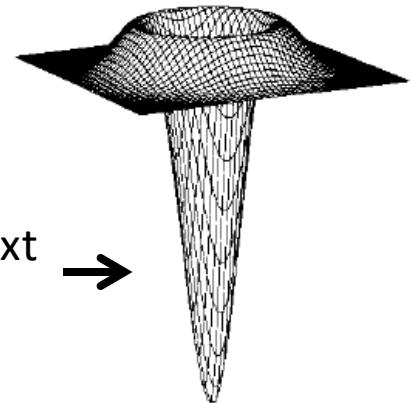
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

$$\nabla^2 h_\sigma(u, v)$$



Laplacian of Gaussian

how does this relate to next
picture?



